

Switching Sides: Choice Structure, Race, and American Support for Foreign Aid

Gavin M. Medina*

September 28, 2026

Abstract

How do country race and choice structure shape white Americans' support for foreign aid? I argue that omnibus questions collapse varied components into a single aggregate evaluation, making conditional opposition appear categorical. In a pre-registered survey experiment, non-Hispanic white Americans evaluated a \$50 million omnibus aid proposal for a Black-majority or white-majority country and then considered the same seven programs separately. Only 36% support the omnibus proposal, but 89% support at least one component, including 76% of omnibus opponents. Respondents share a common preference ordering of programs and differ primarily in how far down they extend support. The country-race cue shifts these funding thresholds conditional on ethnocentrism rather than reorganizing the hierarchy. Moreover, the race-by-ethnocentrism interaction is substantially larger for the omnibus than for the itemized programs. Omnibus measures therefore conceal a coalition of conditional supporters while displaying sharper ethnoracial division than the underlying individual component choices reveal.

Keywords: foreign aid, race, mass politics, public opinion, survey experiments

*Ph.D. Candidate, Department of Politics, Princeton University. gavinmm@princeton.edu

I thank Tali Mendelberg, Christopher Blair, Ismail White, Bocar Ba, and Darius Cozart for helpful feedback. Princeton University's Center for the Study of Democratic Politics provided financial support. Princeton University's Institutional Review Board approved the human subjects research involved in this study.

1 Introduction

Americans are conventionally characterized as reluctant to support foreign aid. For decades, surveys have found that many Americans believe the United States spends too much on aid and favor reducing it (Gilens 2001; Chicago Council on Global Affairs 1974–2024; Scotto et al. 2017; Kates et al. 2019). Scholars attribute this opposition to exaggerated estimates of the foreign aid budget, limited internationalism, partisan polarization, preferences for low government spending, and negative attitudes toward aid recipients (Gilens 2001; Milner and Tingley 2013, 2015; Prather 2024). Overall, studies largely support the view that foreign aid is broadly unpopular among Americans.

Yet this view generally treats foreign aid as a single, race-neutral category of government spending. Existing research typically measures foreign aid opinion through either general questions that treat aid as one policy or questions that examine a single sector, such as humanitarian, military, or democracy assistance, in isolation (Baker 2015). Foreign aid, however, encompasses programs as different as medicine, food, transportation infrastructure, democracy promotion, military equipment, and direct cash transfers to citizens or governments. These programs differ in their purposes, perceived effectiveness, and susceptibility to misuse (Milner and Tingley 2013, 2015). General questions therefore combine distinct programs into a single evaluation and lack a common meaning because respondents vary in what they understand “foreign aid” to include (Williamson 2019). Existing evidence shows that respondents favor some broad categories of aid over others (Doherty et al. 2020), but it does not determine whether general opponents would affirmatively support any particular program.

White Americans provide a demanding case in which to reconsider the conventional account. Recent research finds that Black Americans are more supportive than white Americans of foreign aid, combating world hunger, and improving living standards abroad (Green-Riley et al. 2026). If substantial program-specific support emerges even among a population that is comparatively skeptical of foreign assistance, then apparent opposition to aid is

unlikely to be uniformly categorical. This raises a first question: are white Americans categorically opposed to foreign assistance, or are they opposed to particular ways of providing it?

White Americans also provide a theoretically relevant case for examining how recipient-country race structures foreign aid preferences. Prior research demonstrates that recipient race can alter white Americans' evaluations of humanitarian assistance (Baker 2015). Whether recipient-country race produces a common response among white Americans, however, or divides them according to their ethnoracial attitudes remains unclear. Ethnocentrism may be especially relevant because it captures the extent to which individuals privilege their ethnic group and its interests when evaluating other groups (Bizumic and Duckitt 2012; Bizumic, Monaghan, and Priest 2021). This raises a second question: does country race uniformly increase or decrease support, or does its effect depend on respondents' ethnoracial attitudes?

These ethnoracial divisions also provide a demanding test of the stability of aid preferences. Country race and ethnocentrism could affect not only how much assistance respondents support but also which programs they prioritize. If these divisions produce competing preferences over aid instruments, the ordering of the programs should vary across country-race conditions and levels of ethnocentrism. A common hierarchy that persists even when the country-race cue produces opposing responses across ethnocentrism would instead indicate that disagreement concerns how far respondents extend assistance more than which programs they consider relatively most desirable.

I argue that evaluating foreign aid as a single policy object can make conditional opposition appear categorical. A general question records whether a respondent supports foreign aid overall but cannot show which of its possible uses the respondent accepts. A respondent who rejects every form of aid and one who supports medicine and food while opposing foreign aid in general can therefore give the same omnibus response. Program-level choices reveal both the additional support and the structure of preferences concealed by that aggregate

evaluation.

I further argue that averaging across respondents can conceal the influence of recipient race. Ethnocentrism should condition how white Americans respond to the racial composition of a recipient country. In particular, the effect of the Black-majority cue should become less favorable as ethnocentrism increases. The cross-cutting influence of ethnocentrism should result in offsetting responses, producing small average effects even when race meaningfully structures support. Crucially, this heterogeneity does not necessitate divergent program preferences. If respondents broadly agree about the relative desirability of aid instruments, recipient race may instead move their acceptance thresholds along a stable hierarchy. Respondents lower in ethnocentrism may extend assistance farther down that hierarchy for a Black-majority country relative to the white-majority, whereas those higher in ethnocentrism may stop funding sooner. An omnibus evaluation compresses these different stopping points into a single response. The race-by-ethnocentrism relationship may therefore appear more pronounced in the omnibus evaluation than across the program choices, where respondents can express their shared support for programs near the top of the hierarchy. This threshold account therefore predicts racial polarization in the extent of support without requiring competing rankings of the programs themselves.

I test this argument with a preregistered survey experiment among non-Hispanic white Americans.¹ The sample was quota-balanced by political ideology to ensure sufficient variation for the preregistered moderation tests. Respondents evaluated a \$50 million foreign aid package containing seven programs: medicine, food, transportation infrastructure, democratic elections, military equipment, cash transfers to citizens, and cash transfers to the recipient government. In the omnibus-plus-itemized arm, respondents first evaluated the proposal generally and then decided whether to fund each program separately (Glaser 2002). A second arm received only the itemized questions, allowing me to assess whether program-level support is the result of seeing the omnibus item first. The experiment also randomized

¹This paper centers on white Americans who do not identify as Hispanic, although white Hispanics are a meaningful subgroup in their own right.

a bundled country-race cue: respondents saw either a Black leader of a Black-majority country or a white leader of a white-majority country. I refer to these as the Black-country and white-country conditions.

The omnibus response conceals substantial conditional support. Only 36% of respondents support foreign aid on the general item, whereas 89% support at least one of the seven named programs. Among those who oppose foreign aid generally, 76% nevertheless fund at least one program when given separate choices. I call these respondents *switchers* because they move from opposition to the aggregate policy category to support for at least one specific program. Because both questions identify the same set of programs, the comparison holds the named policy content constant while changing the level at which respondents express their preferences.

The program-level choices also reveal a common ordering. Medicine and food occupy the top of the hierarchy, transportation infrastructure and democratic assistance fall in the middle, and military equipment and cash transfers occupy the bottom. Nearly two-thirds of respondents make choices that exactly follow this ordering with a personal funding threshold.² Respondents therefore differ less in how they rank aid instruments than in how far down the program hierarchy they are willing to extend support.

I find that recipient race operates differently across levels of choice. The Black country cue increases omnibus support by 11.4 percentage points but changes average itemized-program support by only 1.1 points. These small average program-level effects conceal opposing responses across levels of ethnocentrism: each one-standard-deviation increase in ethnocentrism makes the cue's effect on pooled program support 4.3 percentage points more negative. The resulting division is substantially sharper in the omnibus evaluation, where the race-by-ethnocentrism interaction is 5.0 percentage points more negative than in the program decisions.

²Overall, 95.4% of respondents fall within one choice of the common ordering. Among respondents who fund some but not all programs, the cases most informative about the ordering, 58.2% share the common ordering.

Yet recipient race does not appear to reorganize the aid hierarchy. The race-by-ethnocentrism interactions have the same negative direction across all seven programs and do not differ detectably from one another. Moreover, the aggregate program rankings remain nearly identical across the four combinations of recipient race and ethnocentrism, with rank correlations ranging from .93 to 1.00. The stability of this hierarchy is especially notable because it persists across the same ethnoracial divisions that change respondents' overall willingness to provide assistance. Even under a country-race cue whose effect becomes less favorable as ethnocentrism increases, respondents continue to prioritize essentially the same programs. Ethnoracial attitudes divide respondents over how far to extend assistance, not principally over which forms of assistance are most desirable. Overall, these results suggest that the country-race cue primarily moves respondents along a largely shared program hierarchy rather than reorganizing the hierarchy itself.

These findings make three contributions. First, they qualify the conventional characterization of foreign aid as broadly unpopular. Even among non-Hispanic white Americans—a population that prior research identifies as comparatively skeptical of foreign assistance—omnibus opposition conceals extensive support for particular programs. Second, the results uncover a durable structure beneath this apparent opposition. Respondents largely share a hierarchy of aid instruments while differing in how far down it they are willing to extend support, and that hierarchy remains stable across country-race conditions and levels of ethnocentrism. Third, the findings show that measurement level changes the ethnoracial division researchers observe. The omnibus evaluation displays substantially stronger ethnocentric moderation than the component decisions, while average treatment effects conceal the opposing reactions producing that moderation. General measures can therefore make foreign aid opinion appear both more categorically opposed and more ethnoracially divided than the underlying program choices reveal.

2 Foreign Aid Opinion, Country Race, and Program Heterogeneity

Prior work shows that opposition to foreign aid partly reflects misperceptions about the size of the aid budget. Americans substantially overestimate foreign aid spending, and providing accurate information can increase support (Gilens 2001; Hurst et al. 2017; Scotto et al. 2017). Yet budget size is only one source of ambiguity. General foreign aid questions also invite respondents to infer what assistance includes and how it will be delivered (Williamson 2019). A respondent may oppose an unspecified aid package while supporting medicine, food, or other programs when they are evaluated separately. Opposition may therefore appear broader not only because respondents overestimate spending but also because the general category combines different policies.

Recipient identity also shapes how Americans evaluate assistance. Baker (2015) finds that white Americans are more supportive of humanitarian aid to Black than white recipients when assistance is monitored by white aid workers. He further finds that the more favorable response to Black recipients is concentrated in in-kind humanitarian assistance rather than cash transfers. This finding suggests that the effect of recipient race may depend on the particular form of assistance respondents evaluate. Other existing research generally examines broad support, individual aid domains, or respondents' relative priorities among aid packages (Doherty et al. 2020). It has not connected the same respondent's general evaluation to separate funding decisions spanning humanitarian, military, democracy, infrastructure, and cash assistance. Further, research on support for particular aid programs and issue areas has largely neglected race and racial attitudes.

By comparing an omnibus evaluation with separate funding decisions across seven programs, this paper distinguishes general opposition from opposition to particular forms of assistance. The design connects research on foreign aid opinion and recipient racial identity to the broader question of how policy aggregation structures expressed support. It also

shows whether country race changes the relative ordering of aid programs or instead moves respondents along a hierarchy that remains largely shared.

3 Choice Structures: Omnibus and Itemized Support

Choice structure matters not only because it affects measured support, but also because it determines which political coalitions are visible and which combinations of policies can attract support. How questions are asked can shape reported opinion (Chong and Druckman 2007). In the domestic spending domain, Jacoby (2000) finds that Americans express greater support when spending is presented program by program rather than in general terms. Riker's (1986) concept of heresthetics provides the broader theoretical logic: political outcomes can change when the dimensions or alternatives available in a choice are restructured.

Glaser (2002) provides the closest empirical precedent for the relationship between choice structure and race. In Jackson, Mississippi, a checklist ballot allowed voters to select which individual school projects should receive funding rather than requiring them to accept or reject the complete bond package. This choice structure altered the coalition willing to fund projects for majority-Black schools, a pattern Glaser also reproduces in a survey experiment. Because the omnibus proposal identified the same potential projects, the checklist did not merely provide new information about how the funds would be used. Instead, it allowed voters to distinguish among projects that the all-or-nothing vote bundled together.

Existing studies therefore show that itemization can change both the level of support and the coalition willing to fund projects benefiting minority groups. What they do not establish is the preference structure that produces this divergence. Respondents who oppose an omnibus proposal may reject every component, or they may support some components while opposing others. Moreover, citizens may disagree about the relative desirability of the components, or they may share a common ordering but differ in how far down that ordering

they are willing to provide support. This paper uses the comparison between omnibus and itemized foreign aid to distinguish among these possibilities.

Foreign aid is an especially good candidate to examine this difference. Foreign aid omnibus questions combine programs that differ substantially in their purposes, implementation, and perceived desirability. Opposition to an omnibus question therefore does not establish that a respondent categorically opposes every form of assistance it contains.

The comparison between omnibus and itemized aid can reveal three different structures of foreign aid opinion. The first is *categorical opposition*. Respondents who oppose the omnibus package may reject foreign assistance as a policy domain. If so, they should continue to reject every program when the package is disaggregated. Their omnibus response would accurately represent opposition to all forms of aid contained in the proposal.

The second possibility is *conditional opposition with competing program orderings*. Some omnibus opponents may support individual programs, but respondents may disagree about which programs are most desirable. One respondent might favor military assistance over humanitarian aid, while another might favor cash transfers over democracy promotion. Under this structure, itemization would reveal additional support, but that support would not conform to a common hierarchy. The disagreement observed at the package level would reflect both different levels of support and competing evaluations of the programs themselves.

The third possibility is a *common ordering with heterogeneous acceptance thresholds*. Respondents may largely agree about the relative desirability of the programs but differ in how far down that ordering they remain willing to provide funding. A respondent with a high acceptance threshold might support only the most preferred program. A respondent with a lower threshold might support several programs, while another might support the entire package. These respondents would disagree about how many programs the United States should fund without necessarily disagreeing about which programs are relatively more desirable.

The itemized checklist distinguishes categorical opponents from conditional supporters.

I call respondents who oppose the omnibus package but support at least one itemized program “switchers.” The term describes movement from opposition to the complete bundle to support for at least one component; it does not imply that itemization changed their underlying preferences. A substantial population of switchers would indicate that omnibus measures classify some conditional supporters as opponents of foreign aid generally.

Hypothesis 1: *Switching: Omnibus opposition overstates opposition to all types of aid. Presented with an itemized list of programs, a portion of omnibus opponents will support at least one program.*

The programs that switchers support provide evidence about the substantive ordering of foreign aid instruments. Concrete humanitarian programs should occupy the top of that ordering. Medicine and food address immediate and visible needs, and respondents can more readily connect them to identifiable benefits. Cash assistance should occupy the bottom, particularly cash transfers to the recipient government. Cash provides recipients with greater discretion, makes the ultimate use of funds less visible to donors, and may evoke concerns about corruption or misuse. Military assistance should also receive comparatively limited support because it does not directly address humanitarian need and may implicate respondents in coercive action.

This expected hierarchy does not require respondents to support all humanitarian assistance or reject all cash assistance. It predicts that medicine and food will receive more support than programs with less immediate or less visible benefits. If switchers are conditional rather than categorical opponents, their support should be concentrated near the top of this ordering.

Hypothesis 2: *Switching is concentrated in concrete humanitarian programs. Among omnibus opponents who support at least one itemized program, support is highest for medicine and food and lowest for cash assistance, particularly cash transfers to the recipient government.*

A separate allocation task provides a second way to observe the ordering. Binary funding choices show whether each program clears a respondent’s acceptance threshold, but they do not measure the relative priority respondents assign to the programs they consider. Requiring respondents to divide a fixed budget across the seven programs forces a direct tradeoff. If switchers’ conditional support is concentrated in humanitarian assistance, they should devote a larger share of their budgets to medicine and food and relatively little to cash transfers. This concentration should be more pronounced among switchers than among respondents who support the complete omnibus package.

Hypothesis 3: *Relative to omnibus supporters, switchers allocate a larger share of the aid budget to humanitarian programs than to cash assistance.*

H1-H3 establish whether omnibus opposition conceals conditional support and where that support is concentrated. The program-level choices also permit a stronger test of preference structure. If respondents possess competing orderings, their choices should not consistently follow a single program hierarchy. If they share an ordering but apply different acceptance thresholds, their choices should form a cumulative pattern. I therefore evaluate the proportion of respondents whose choices conform to the modal ordering, the number of deviations required to make each choice set cumulative, and the scalability of the seven programs.

The omnibus and itemized outcomes should not be interpreted as interchangeable measures or as identifying a causal effect of itemization. Supporting an omnibus package and supporting at least one component are substantively different choices. Their comparison instead identifies preferences that the omnibus measure cannot reveal. The omnibus item explicitly named all seven programs that subsequently appeared in the itemized checklist. The comparison therefore holds the package’s stated contents constant while changing the level at which respondents express support. A separate experimental arm receives only the itemized checklist, allowing the analysis to test whether program-level support is an artifact

of having first evaluated the omnibus package.³ The design therefore distinguishes conditional support from categorical opposition while separately assessing whether the order of the questions changes itemized choices.

4 Recipient Race and Support for Foreign Aid

A bundled country-race cue may produce different responses among white Americans with different ethnoracial attitudes. By jointly varying the apparent race of the leader and the stated racial majority of the recipient country, the experiment tests whether a small average treatment effect combines opposing responses across levels of ethnocentrism. The program-level choices then show whether this heterogeneity changes respondents' relative priorities or instead moves their funding thresholds along a stable hierarchy.

The preregistration also specified paternalism, perceived agency, and perceived living standards as possible pathways linking recipient race to aid preferences. Because these measures were collected after the funding outcomes, I treat them as descriptive pathway diagnostics rather than identified mediators and report the complete analyses in the appendix.

4.1 Ethnocentric In-Group Preference: Ethnocentrism and Racial Resentment

Ethnocentrism, as formulated by Bizumic et al. (2021), is a general disposition to privilege one's own ethnic group over others. Because the measure is defined without reference to a particular out-group, it can be applied to evaluations of foreign populations.⁴ This general construction is useful in an international setting in which respondents evaluate the population of a foreign country rather than a domestic racial group.

³Support for specific programs is statistically indistinguishable across arms. The estimated difference is roughly two percentage points (see Table A6), and the order-by-race interaction is also null. These results provide evidence that program-level support is not an artifact of question order.

⁴See Kinder and Kam (2009) for an alternative operationalization.

Recent work shows that race structures foreign policy attitudes across domains. Black and white Americans differ in their views of foreign policy (Green-Riley et al. 2026), while white identity and racialized cues shape perceptions of foreign adversaries and support for trade partners (Song 2026; Mutz, Mansfield, and Kim 2021).⁵ Whether ethnocentrism similarly conditions white Americans’ responses to country-race cues in foreign assistance remains unclear.

If ethnocentrism produces in-group preference, the effect of assignment to the Black-country condition, relative to the white-country condition, should become less favorable as ethnocentrism increases. White respondents higher in ethnocentrism should extend more favorable treatment when presented with the white-country condition. At lower levels of ethnocentrism, this preference should weaken and may reverse because the lower end of the scale captures greater concern for the welfare of ethnoracial out-groups. Consequently, a small average effect of the country-race cue would not necessarily indicate that race is unimportant. It could instead combine more favorable responses to the Black-country recipient among respondents lower in ethnocentrism with less favorable responses among those higher in ethnocentrism.

Hypothesis 4: *Ethnocentrism moderates the effect of the country-race cue on switching and cash-program support, such that the effect of the Black-country cue becomes less favorable as ethnocentrism increases.*

Beyond this preregistered hypothesis, choice structure may condition the magnitude of the ethnocentric division. An omnibus question requires respondents to register a single evaluation of the complete proposal, whereas itemized choices allow them to distinguish among aid programs (Glaser 2002). The race-by-ethnocentrism interaction may therefore be more pronounced for the omnibus evaluation than for the program choices. Because these questions measure substantively different outcomes, I treat this cross-outcome comparison

⁵Related conjoint evidence suggests that American support for Asian and Indigenous victims of aggression may also differ across levels of ethnocentrism (Inouye 2026).

as exploratory rather than as an identified causal effect of itemization.

Racial resentment provides a more target-specific alternative moderator. Its designers conceptualized it as a combination of anti-Black affect and traditional American values, including individualism, self-reliance, and the belief that advancement should follow effort (Kinder and Sears 1981). The items, however, refer specifically to Black Americans and American racial history. A battery concerning whether Black Americans could improve their circumstances through greater effort may therefore be less portable to evaluations of a foreign government than a target-general measure of ethnocentrism.⁶ If racial resentment captures a broader anti-Black orientation, it should moderate the effect of the Black-country cue in the same direction as ethnocentrism. Because its content is tied more closely to the domestic context, I expect this interaction to be weaker.

Hypothesis 5: *Racial resentment moderates the effect of the country-race cue on switching and cash-program support in the same direction as ethnocentrism (H4), but with a weaker interaction.*

4.2 Preregistered Pathways: Paternalism, Agency, and Perceived Need

The preregistration identified racial paternalism as a potential pathway connecting recipient-country race to aid preferences. Paternalism combines a willingness to assist with the belief that recipients require external guidance or control (Jackman 1994; Moore 2020). Applied to the present setting, the argument predicts that assignment to the Black-country condition will increase beliefs that poor countries require American guidance and that the United States should constrain recipient-government choices. Those beliefs should, in turn, be associated with opposition to cash assistance and with switching from opposition on the omnibus item to support for selected programs. I measure the benevolence and control dimensions

⁶Parajon (2025) finds that racial resentment predicts climate attitudes.

separately and through a composite paternalism index.⁷

I also reproduce Baker’s (2015) measures of perceived agency and perceived living standards as alternative pathways. The agency measure assesses whether the country-race cue changes respondents’ perceptions of recipients’ capacity to improve their circumstances. The living-standards measure assesses whether respondents infer greater material need in the Black-country condition even though the vignette holds national income constant. Comparing these measures with paternalism allows the preregistered analysis to distinguish beliefs about recipient guidance, agency, and material need.

Hypothesis 6: *Racial paternalism mediates race-conditioned switching. Assignment to the Black-country condition increases reported paternalism, which is associated with switching and opposition to cash assistance. The mediated relationship is expected to be strongest for cash transfers to the recipient government and absent for military assistance; benevolence is expected to mediate omnibus support; and paternalism is expected to account for more of the country-race relationship than the agency and living-standards measures.*

These candidate mediators were measured after treatment but also after respondents completed the omnibus and itemized funding decisions. The randomized treatment therefore identifies whether the country-race cue changes respondents’ subsequent reports of paternalism, agency, or living standards. It does not establish that those reports preceded or caused their aid choices. The associations may instead reflect rationales respondents reported for decisions they had already made. I therefore treat the indirect-effect estimates as preregistered pathway diagnostics rather than as definitive evidence of causal mediation.

5 Method

To examine how support differs when respondents evaluate foreign aid as a complete omnibus package and as a set of itemized programs, I fielded a preregistered survey exper-

⁷Appendix Section 9.8 reports the exact question wording.

iment.⁸ Respondents first completed demographic questions and pretreatment measures of partisanship, political ideology, internationalism, self-monitoring, ethnocentrism, and racial resentment.⁹ The treatment jointly varied the leader’s apparent race and the country’s stated racial majority: respondents saw either a Black leader of a Black-majority country or a white leader of a white-majority country. The experiment therefore identifies the effect of a bundled recipient-race cue rather than the separate effect of leader race or population composition.¹⁰

Main-arm respondents then completed two principal funding decisions. They first evaluated the complete \$50 million aid package on an omnibus support-or-oppose question, on a 5-point Likert scale, and then indicated whether they supported funding each of its seven programs separately. The omnibus proposal explicitly named the same seven programs that subsequently appeared in the itemized checklist. Respondents therefore evaluated the complete seven-program package before identifying which of its constituent programs they supported (Glaser 2002). A separate randomly assigned arm received only the itemized checklist, allowing me to test whether program-level support depended on having first evaluated the omnibus package.¹¹

After the funding decisions, respondents allocated the \$50 million across the seven programs.¹² This task provides a complementary measure of the relative priority assigned to each program. Respondents then ranked cash, supplies, and training as delivery formats for each of the five non-cash programs. This task tests whether reluctance toward cash extends beyond the two programs explicitly defined as cash assistance. Respondents subsequently completed an open-ended question about the considerations underlying their choices. Embedded JavaScript recorded paste events and typing speed to implement the preregistered

⁸The preregistration is available at OSF.

⁹Recent experimental evidence finds no systematic evidence that measuring moderators before treatment alters treatment effects (Sheagley and Clifford 2025).

¹⁰The cue worked as intended: 99.3% of respondents in the Black-majority condition and 97.5% in the white-majority condition correctly identified the recipient country’s racial composition (Table A4). The leader images were generated using ChatGPT.

¹¹Program order was randomized at the respondent level and carried through each subsequent battery.

¹²A pie chart helped respondents visualize their allocations.

screen for potentially AI-assisted responses (Westwood 2025; Westwood and Frederick 2026).

Respondents next saw a reminder of their treatment condition and completed the post-treatment mediator items in randomized order. Two measures replicate Baker’s (2015) perceived-agency and perceived-living-standards measures. The third is a paternalism index composed of benevolence and control subscales. Whereas Baker’s agency measure assesses the perceived capacity of poor foreigners to improve their circumstances, my measure assesses whether poor countries require American guidance and whether the United States should override or constrain recipient-government choices. After completing the mediator battery, respondents answered a separate set of questions about the recipient country’s leader and government that is reserved for other work.

Finally, respondents evaluated each of the seven programs along six dimensions: whether it reduces theft, is easy to verify, benefits the U.S. economy, increases U.S. influence, saves lives, and is morally right. They then reported the religion and region they associated with the recipient country. Because the program evaluations followed the funding choices and other post-treatment questions, I interpret them as subsequent reported evaluations rather than beliefs that necessarily caused respondents’ funding decisions.

5.1 Sample

The survey was fielded on August 9, 2026 on Prolific, restricted to non-Hispanic white Americans, with recruitment quotas on self-reported ideology (35 percent liberal, 30 percent moderate, 35 percent conservative) to ensure adequate power for the ideological and racial attitude moderation tests.¹³ Respondents who failed both of two attention checks were screened out before treatment. The final data contain 2,026 treated respondents; the preregistered keystroke rule — a paste event or typing above 800 characters per minute on the open-ended item — excludes 30, leaving 1,996: 1,769 in the main arm and 227 in the itemized-only arm. All analyses use only completed responses, as preregistered.

¹³For descriptives of the sample see Table A2.

All multi-item scales perform adequately (Table A5 in the appendix): ethnocentrism $\alpha = .77$ and racial resentment $\alpha = .94$ (both pre-treatment); the post-treatment mediators — the four-item paternalism index ($\alpha = .74$; its benevolence and control subscales at .70 and .69, correlated at $r = .46$), Baker’s perceived-agency scale ($\alpha = .74$), and perceived living standards ($\alpha = .78$). Covariates in all adjusted models are ideology (mean-centered), Independent and Republican indicators, internationalism, self-monitoring, evangelical identification, age, gender, education, and income.

6 Analysis Strategy

The analysis proceeds in three stages. First, I compare support for the complete aid proposal with support for its seven constituent programs. I identify respondents who oppose the omnibus proposal but support at least one program as switchers and use the itemized-only arm to assess whether program-level support depends on prior exposure to the omnibus evaluation.

Second, I examine the structure of the itemized choices. I compare program support and budget allocations across omnibus-response groups and assess whether respondents’ choices conform to a common program ordering with heterogeneous funding thresholds. The cumulative-scale analyses are exploratory and were not specified in the preregistration.

Third, I estimate the average and conditional effects of the bundled country-race cue. The preregistered tests examine moderation by ethnocentrism and racial resentment for switching and cash-program support. Exploratory analyses extend these tests to pooled program support, the number of programs funded, ordering stability, and the difference between moderation in the omnibus and program-level outcomes. Respondent-level models use HC2 robust standard errors, and models containing repeated program observations cluster standard errors by respondent.

The appendix reports the remaining preregistered delivery-mode and pathway analyses,

along with exploratory analyses of respondents’ subsequent program evaluations. Because the candidate mediators and program evaluations were measured after the funding decisions, I interpret them as descriptive pathway diagnostics rather than identified causes of respondents’ choices.

7 Results


7.1 Omnibus Opposition Conceals Conditional Support

Support for at least one component of the aid proposal substantially exceeds support for the complete proposal presented as a single omnibus evaluation. In the main arm ($n = 1,769$), 36.4% of respondents support the omnibus proposal, whereas 89.0% support at least one of its seven constituent programs. The within-respondent difference is 52.6 percentage points ($p < .001$), confirming Hypothesis 1. Five of the seven programs individually receive more support than the complete proposal (Figure 1, panel a; Appendix Table A7). This comparison holds the proposal’s named contents constant, but it contrasts support for the complete set of programs with support for any individual component. It therefore reveals conditional support rather than identifying a causal effect of itemization.

The high level of itemized support is not explained by a substantively large question-order effect. The preregistered between-arm test compares itemized support in the main arm, which answered the omnibus item first, with support in the itemized-only arm, which never saw it. The estimated order effect is -0.023 ($SE = 0.018$, $p = .20$) and does not differ by recipient race ($p = .93$; Appendix Table A6). The itemized-only arm also reproduces the same broad program hierarchy: medicine and food receive the most support, infrastructure and elections occupy the middle, and military equipment and cash transfers receive the least support. A joint test cannot reject the null that recipient race leaves support for all seven programs unchanged in this arm ($\chi^2(7) = 9.65$, $p = .21$; Appendix Table A9).

Among respondents who oppose the omnibus proposal ($n = 811$), 76.4% support at

Figure 1: Program Support and Budget Allocations by Omnibus-Response Group



Notes: Panel (a) shows the share of each group supporting funding for the named program on the itemized checklist; the dashed line denotes the share of all main-arm respondents supporting the complete \$50 million proposal on the omnibus item (36.4%). Panel (b) shows the mean percentage of the \$50 million budget that each group assigns to the program in the allocation task. Omnibus supporters (black circles, $n = 644$) support the complete proposal; soft switchers (blue squares, $n = 311$) neither support nor oppose the proposal and fund at least one program; switchers (crimson triangles, $n = 620$) oppose the proposal but fund at least one program. Thick and thin whiskers denote 90% and 95% confidence intervals. Exact values appear in Appendix Table A10; covariate-adjusted estimates appear in Appendix Figure A1. Data: Prolific, August 2026.

least one individual program. These 620 respondents constitute the switchers. Among the 314 respondents who neither support nor oppose the omnibus proposal, 99.0% fund at least one program. Approximately three-quarters of omnibus opponents are therefore not categorically opposed to every form of foreign assistance. They reject the proposal as a whole while remaining willing to fund at least one of its components.

The full omnibus response scale reveals the same pattern more generally (Table 1). The mean number of programs funded rises steadily from 1.44 among respondents who oppose the omnibus proposal to 5.19 among those who support it. Yet the correspondence is not categorical: 56.3% of respondents selecting “oppose” fund at least one program, while only 16.4% of those selecting “support” fund all seven. Omnibus responses therefore track how far respondents extend assistance while concealing substantial variation in their program-level

funding thresholds.

Table 1: Omnibus responses and the number of programs funded

Omnibus response	Programs funded (row %)								Mean
	0	1	2	3	4	5	6	7	
Oppose	43.7	8.9	23.5	12.4	7.5	3.8	0.0	0.3	1.44
Oppose somewhat	6.6	2.3	26.4	29.3	22.0	11.1	1.8	0.5	3.01
Neither	1.0	1.0	7.6	16.6	37.3	26.4	5.7	4.5	4.13
Support somewhat	0.0	0.0	2.9	9.9	32.2	38.1	13.0	3.9	4.60
Support	0.0	0.0	1.3	4.4	20.8	37.7	19.5	16.4	5.19

Notes: Each row shows the distribution of the number of itemized programs funded among main-arm respondents giving the stated omnibus response, in row percentages, with the row mean in the final column. Percentages may not sum to 100 because of rounding. Main arm, $N = 1,769$. Data: Prolific, August 2026.

The difference between omnibus and itemized support consequently reveals more than additional support for particular programs. The omnibus evaluation makes respondents who reject every component observationally equivalent to those who accept some programs but not the complete proposal. The next question is whether these conditional supporters hold idiosyncratic preferences over the programs or whether their choices follow a common structure.

7.2 A Common Hierarchy with Heterogeneous Funding Thresholds

The itemized choices reveal substantial agreement about the relative desirability of the seven programs. Medicine and food occupy the top of the hierarchy, transportation infrastructure and democratic elections occupy the middle, and military equipment and cash transfers occupy the bottom. Omnibus supporters, soft switchers, and switchers differ primarily in how far down this hierarchy they extend support rather than in which programs they prioritize.

The program-level results support Hypothesis 2. As panel (a) of Figure 1 shows, 94.0% of switchers support medicine and 93.1% support food, compared with only 3.4% who sup-

port cash transfers to the recipient government. Among switchers, the mean support rate for medicine and food exceeds the mean support rate for the two cash programs by 82.6 percentage points ($p < .001$). Conditional support is therefore concentrated among the programs located at the top of the common hierarchy.

The allocation results provide complementary evidence and support Hypothesis 3. As panel (b) shows, switchers allocate 74.5% of the \$50 million budget to medicine and food, compared with 56.9% among omnibus supporters. Switchers allocate only 6.1% to the two cash programs, compared with 13.9% among omnibus supporters. The allocation task therefore reproduces the hierarchy observed in the binary funding choices: respondents who reject the complete proposal assign their available resources disproportionately to its humanitarian components.

The hierarchy is not unique to switchers. Omnibus supporters fund most of the programs, whereas switchers and soft switchers concentrate their support near the top of the hierarchy. Respondents whom the itemized choices identify as categorical opponents reject all seven programs. These groups therefore differ principally in their funding thresholds, not in the broad ordering they assign to the programs.

An exploratory non-preregistered cumulative-scale analysis provides additional evidence of this structure. The same structure appears at the individual level. Among respondents who fund some but not all programs—the cases that directly identify an ordering—58.2% follow the modal hierarchy exactly. Across the full sample, 64% fund every program above a personal cutoff in that ordering and reject every program below it. If respondents selected the programs independently at their observed marginal rates, only 36% would produce such a cumulative pattern. An additional 31% deviate from the modal ordering on only one program, meaning that 95.4% of respondents fall within one choice of an ordering-plus-threshold pattern. The seven programs also form a strong cumulative scale (coefficient of reproducibility = .942; Loevinger’s $H = .615$; Appendix Section 3.3). These results are consistent with a common ordering accompanied by heterogeneous funding thresholds

(Guttman 1950; Mokken 1971; van Schuur 2003).

A separate delivery-mode ranking task provides additional evidence that the low position of cash assistance extends beyond the two programs explicitly defined as cash transfers. Respondents rank cash last among training, cash, and supplies for four of the five non-cash programs and rank it behind supplies for all five (Appendix Table A16).¹⁴ Because these rankings are ancillary to the principal choice outcomes, the complete preregistered results appear in the appendix.

Respondents' subsequent evaluations of the programs also correspond to the hierarchy: medicine and food receive the most favorable evaluations, whereas military equipment and cash assistance receive the least favorable evaluations. Because these evaluations were reported after the funding decisions, I treat them as descriptive characterizations rather than explanations of the program choices and report the complete analysis in Appendix Section 4.

These results do not imply that respondents agree about how much foreign aid the United States should provide or that every respondent possesses exactly the same ranking. They show that disagreement is structured more by heterogeneous funding thresholds than by competing orderings of aid programs. Respondents largely agree about the relative desirability of the seven programs but differ in how many they are willing to fund. The omnibus evaluation records these threshold differences as apparent opposition to foreign aid as a whole.

Itemization therefore reveals a structured disagreement hidden by the omnibus evaluation. This structure permits a more precise analysis of recipient race: the country-race cue may reorganize the program hierarchy itself, or it may move respondents' funding thresholds along a hierarchy that remains largely shared.

¹⁴Food assistance is the only exception to cash ranking last: respondents rank cash slightly ahead of training, with mean ranks of 2.26 and 2.38, respectively. Supplies remain the most preferred delivery mechanism for food assistance.

7.3 Country Race Moves Funding Thresholds Conditional on Ethnocentrism

7.3.1 Average Effects of the Country-Race Cue

I begin with non-preregistered exploratory estimates of the average effects of the randomized country-race cue. I first estimate the average effect of the country-race cue on support for the omnibus aid proposal. Respondents in the Black-country condition support the omnibus proposal at a rate of 42.0%, compared with 30.6% in the white-country condition. The resulting 11.4-percentage-point difference is statistically significant ($p < .001$) and remains approximately unchanged after adjustment for the preregistered covariates (Appendix Table A18). This result is consistent with Baker’s (2015) finding that white Americans extend greater aid support to Black than white recipients.

The average effects on the itemized outcomes are substantially smaller, although the evidence for smallness is stronger at some levels of aggregation than at others. Respondents are approximately equally likely to fund at least one program in the two country-race conditions, and the cue changes average program support by only 1.1 percentage points. Equivalence tests reject effects outside three percentage points for funding at least one program and outside five percentage points for average program support (Appendix Table A18).

The individual-program estimates are similarly small but less precise. The largest differences are approximately four percentage points: +3.9 points for transportation infrastructure and −3.8 points for military equipment. Two estimates reach conventional statistical significance before adjustment for multiple comparisons and point in opposing directions: assignment to the Black-country condition reduces support for military equipment by 3.8 points ($p = .028$) and increases support for cash transfers to the recipient government by 3.3 points ($p = .041$). Neither survives a Benjamini–Hochberg correction across the seven programs. I therefore interpret the aggregate effects as small where they can be bounded and the individual-program estimates as imprecise rather than as established nulls. Appendix

Figure A4 displays the program-specific estimates.

The country-race cue therefore increases omnibus support without producing a comparable average increase in support for the individual programs. One possibility is that the cue operates on the proposal as a whole without affecting particular components. Another is that the program-level averages combine respondents who react to the cue in opposing directions. The moderation analyses evaluate this second possibility.

7.3.2 Preregistered Moderation Tests

Hypothesis 4 predicts that ethnocentrism moderates the effect of the country-race cue on cash-program support and switching. Cash-program support is the respondent’s average support across cash transfers to citizens and cash transfers to the recipient government. The interaction between assignment to the Black-country condition and standardized ethnocentrism is -0.033 ($p = .021$). Each one-standard-deviation increase in ethnocentrism therefore makes the effect of the Black-country cue on cash-program support 3.3 percentage points less favorable, supporting Hypothesis 4 for cash assistance.

The corresponding interaction for switching is -0.055 ($p = .051$). Although directionally consistent with Hypothesis 4, this estimate should be interpreted descriptively because switching is defined among respondents who opposed the omnibus proposal, and omnibus opposition is itself a post-treatment outcome. Omnibus opposition is 48.4% in the white-country condition and 43.4% in the Black-country condition. Moreover, opponents in the Black-country condition score 0.18 points higher in ethnocentrism than opponents in the white-country condition ($p = .001$). Treatment assignment therefore changes the composition of the opponent pool, preventing a clean causal interpretation of the switching interaction.

Hypothesis 5 proposes racial resentment as a more domestically anchored alternative moderator. Its preregistered interactions with the country-race cue are negative but statistically indistinguishable from zero for switching ($p = .38$) and cash-program support ($p = .21$).

Hypothesis 5 therefore receives no clear support.

Because ethnocentrism and racial resentment are correlated ($r = .61$), I also estimate exploratory models containing both moderators and their interactions with the country-race cue. Conditional on racial resentment, the ethnocentrism interaction remains negative and statistically significant for cash-program support (-0.039 , $p = .029$) and pooled program support (-0.048 , $p = .002$), whereas the corresponding racial-resentment interactions are close to zero. The differences between the two interactions are not statistically significant for cash support ($p = .175$) and reach only marginal significance for pooled program support ($p = .058$). These results do not establish that ethnocentrism is the stronger moderator. They instead indicate that ethnocentrism captures variation in program-level responses to country race that is not absorbed by racial resentment.

The omnibus outcome does not exhibit the same distinction. When the two moderators enter simultaneously, their interactions with the country-race cue are both negative and statistically significant, and they are statistically indistinguishable from one another ($p = .86$; Appendix Table A20). Evidence differentiating the target-general and domestically anchored measures is therefore concentrated in the program-level choices rather than in evaluations of the omnibus proposal.

7.3.3 Stable Orderings, Shifting Thresholds, and Stronger Omnibus Moderation

Exploratory analyses show that ethnocentric moderation extends beyond cash assistance and operates primarily through movement along the shared program hierarchy. For omnibus support, the interaction between assignment to the Black-country condition and standardized ethnocentrism is -0.090 ($p < .001$; Appendix Table A19). Each one-standard-deviation increase in ethnocentrism makes the effect of the Black-country cue 9.0 percentage points less favorable. The 11.4-point average advantage for the Black-country recipient is therefore concentrated among respondents lower in ethnocentrism and declines substantially as


ethnocentrism increases.¹⁵

The pattern extends across the seven programs. Each one-standard-deviation increase in ethnocentrism makes the effect of the Black-country cue on pooled program support 4.3 percentage points more negative ($p < .001$; Appendix Table A21). The seven program-specific interactions range from -2.2 to -5.7 percentage points, and all are negative. Five are statistically distinguishable from zero at the .05 level before adjustment for multiple comparisons. Because the individual estimates are exploratory and several are imprecise, the pooled interaction provides the clearest summary of the common pattern.

The conditional racial division is significantly more pronounced in the omnibus evaluation than in the program choices. In a stacked model containing each respondent's omnibus response and seven program decisions, the country-race-by-ethnocentrism interaction is 5.04 percentage points more negative for the omnibus outcome than for the programs ($SE = 0.0187$, $p = .007$). This comparison does not identify a causal effect of itemization because support for the complete proposal and support for an individual program are substantively different outcomes. It nevertheless shows that the omnibus evaluation displays a substantially larger ethnocentric division in responses to country race than the component decisions do.

¹⁵Appendix Section 5.5 relaxes the linear interaction assumption following Hainmueller, Mummolo, and Xu (2019). Binning and kernel estimates produce the same broad pattern of an increasingly negative response to the Black-country cue as ethnocentrism rises, although the flexible estimates indicate some nonlinearity and greater uncertainty at the extremes of the moderator distribution.

Figure 2: Country Race and Support at Lower and Higher Levels of Ethnocentrism



Notes: Points show marginally standardized predicted probabilities of support in the Black-country condition (black circles) and white-country condition (crimson triangles) at one standard deviation below and above mean ethnocentrism. Predictions come from logistic models containing treatment assignment, continuous ethnocentrism, their interaction, and the full covariate set, averaged over the observed covariate distribution. Ethnocentrism is modeled continuously; the displayed values summarize the interaction. Corresponding linear probability models show the same pattern of increasingly negative treatment effects as ethnocentrism rises. The program ordering remains highly stable across panels, while the Black-country advantage at lower ethnocentrism narrows and, for most programs, reverses at higher ethnocentrism. Thick and thin whiskers denote 90% and 95% confidence intervals based on 1,000 draws from the coefficient sampling distribution. Main arm, $N = 1,769$. Data: Prolific, August 2026.

Figure 2 displays the pattern at one standard deviation below and above mean ethnocentrism. At lower ethnocentrism, predicted support is higher in the Black-country condition for each program. At higher ethnocentrism, the estimated differences reverse direction for five of the seven programs and approach zero for medicine and food, although several individual contrasts remain imprecisely estimated. The strongest evidence is therefore not seven independently established reversals but the common negative direction of the interactions and the precisely estimated pooled relationship.

Three additional results support a threshold interpretation. First, a joint test cannot

reject equality of the seven country-race-by-ethnocentrism interactions ($\chi^2(6) = 3.13$, $p = .79$). This test does not establish that the interactions are identical, but their common direction and statistical similarity are consistent with a general shift in support rather than program-specific changes in the effect of country race. An exploratory equivalence test provides affirmative evidence for this interpretation: all seven program-specific deviations from the mean interaction can be bounded within five percentage points, with the largest estimated deviation equal to 2.1 points (intersection-union $p = .041$; Appendix Table A23).

Second, the aggregate ordering of the programs remains highly stable. Ranking the seven programs by support within each combination of country-race condition and ethnocentrism produces Spearman correlations ranging from .93 to 1.00 (Appendix Table A22). Among respondents higher in ethnocentrism, the ordering is identical in the Black-country and white-country conditions. The country-race cue therefore changes levels of support without meaningfully changing which programs receive the most support.

Third, the interaction appears directly in respondents' funding thresholds. In the covariate-adjusted model, each one-standard-deviation increase in ethnocentrism makes the effect of the Black-country cue on the number of programs funded 0.279 programs more negative ($SE = 0.078$, $p < .001$). The cue moves respondents lower in ethnocentrism farther down the shared hierarchy while moving respondents higher in ethnocentrism toward a more restrictive cutoff.

Taken together, these results indicate that country race does not appear to reorganize the aggregate ranking of aid instruments. Instead, it changes how far down a largely shared ordering respondents are willing to extend support, and the direction of that change depends on ethnocentrism. The omnibus evaluation compresses these differences into a single aggregate decision and consequently displays a sharper ethnocentric division than the underlying program choices.


7.4 Racial Polarization Extends to Subsequent Program Evaluations

The same offsetting pattern appears in respondents' subsequent evaluations of what the programs accomplish. The average effect of the country-race cue on a six-item index of affirmative program evaluations is approximately zero (0.1 percentage points, $p = .94$). This average masks opposing responses across ethnocentrism. Each one-standard-deviation increase in ethnocentrism makes the effect of the Black-country cue on the evaluation index 3.8 percentage points more negative ($p < .001$). At one standard deviation below mean ethnocentrism, the cue increases affirmative program evaluations by 4.0 points ($p = .001$); at one standard deviation above the mean, it decreases them by 3.6 points ($p = .012$).

Figure 3 shows that the estimated effect becomes increasingly negative as ethnocentrism rises and crosses zero near the mean of the moderator. The pattern is not confined to one type of evaluation. The joint null that the six country-race-by-ethnocentrism interactions equal zero is rejected ($\chi^2(6) = 17.5$, $p = .008$), and all six interactions are negative. The estimates range from -2.9 percentage points for believing that a program saves lives to -4.3 points for believing that its use is easy to verify, and they do not differ detectably across the six dimensions ($\chi^2(5) = 1.9$, $p = .86$; Appendix Table A24).

The country-race cue therefore produces more affirmative subsequent program evaluations among respondents lower in ethnocentrism and less affirmative evaluations among those higher in ethnocentrism. Averaging across respondents conceals these opposing reactions just as it does for program support. Because the evaluation items followed the funding choices, this correspondence should not be interpreted as evidence that the reported evaluations caused or mediated respondents' program choices. It instead provides convergent evidence that ethnoracial polarization extends beyond funding decisions to evaluations of what identical aid programs accomplish.

Figure 3: Effect of the Country-Race Cue on Subsequent Program Evaluations, by Ethnocentrism



Notes: The line is the estimated effect of assignment to the Black-country condition on a six-item index of affirmative program evaluations across the observed range of standardized ethnocentrism; the shaded band denotes 95% confidence intervals. The index averages respondents' affirmative answers to whether each program saves lives, is morally right, is easy to verify, reduces theft, increases U.S. influence, and benefits the U.S. economy. Estimates come from a stacked model containing one observation per respondent-program, program fixed effects, and standard errors clustered by respondent. The vertical dotted line marks mean ethnocentrism. Rug marks show observed values of ethnocentrism; darker marks indicate greater respondent density. Tabular results appear in Appendix Table A24. Main arm, $N = 1,769$. Data: Prolific, August 2026.

7.5 Preregistered Pathway Diagnostics

The remaining preregistered pathway analyses provide little support for the proposed paternalism mechanism. Assignment to the Black-country condition does not increase composite paternalism or control-oriented paternalism, and the estimated indirect effect through composite paternalism is 0.5 percentage points (95% CI: $-0.3, 1.4$). The preregistered predictions that depend on a treatment-induced increase in paternalism are therefore not supported. Paternalism is associated with which observed omnibus opponents fund selected programs, but this relationship is descriptive because paternalism was measured after the funding decisions and the analysis conditions on a post-treatment outcome. Perceived agency

likewise provides no evidence of an indirect pathway. The country-race cue does lower reported living standards by 3.6 percentage points ($p < .001$) despite identical information about national income, producing estimates consistent with an inferred-need pathway. These estimates are sensitive to modest violations of sequential ignorability, however, and all candidate mediators were measured after respondents made their funding decisions. I therefore treat the pathway results as preregistered diagnostics rather than identified causal mediation. Appendix Section 7 reports the complete hypotheses, model-based indirect effects, preregistered comparisons, and sensitivity analyses.

8 Discussion and Conclusion

Generic policy questions can make conditional opposition appear categorical. Only 36.4 percent of respondents support the omnibus aid package, whereas 89 percent support at least one of its constituent programs. Among respondents who oppose the package, 76.4 percent nevertheless fund at least one program when its components are evaluated separately. These differences should not be interpreted as the causal effect of itemization: supporting a complete package and supporting at least one component are substantively different choices. Their comparison instead reveals preferences that an omnibus evaluation cannot express. Most respondents classified as opponents of foreign aid are not opposed to every form of assistance. Their opposition is conditional on what the package contains.

The program hierarchy also identifies the structure of this conditional opposition. Omnibus opponents could have held competing rankings of the programs, with some favoring humanitarian assistance and others preferring military aid or cash transfers. Instead, respondents largely share a common ordering. Medicine and food occupy the top of that ordering, transportation infrastructure and democratic assistance fall in the middle, and military equipment and cash transfers occupy the bottom. Among respondents who fund some but not all programs—the cases that directly identify an ordering—58.2% follow the

common hierarchy exactly. Across the full sample, 95.4% fall within one choice of that hierarchy, and the programs form a strong cumulative scale. Respondents therefore disagree primarily about how many programs the United States should fund, not about which programs are relatively more desirable. The omnibus format converts these differences in acceptance thresholds into apparent disagreement over foreign aid as a whole.

The central finding is therefore not simply that omnibus opponents support some individual programs. It is that omnibus and itemized choices reveal different representations of the same structured disagreement. Respondents largely agree that medicine and food are more desirable than military equipment and cash transfers, but they differ in where they stop funding along that hierarchy. Recipient race moves those stopping points in opposing directions across levels of ethnocentrism without substantially changing the ordering itself. The omnibus evaluation compresses these differences into a single decision and displays a race-by-ethnocentrism interaction that is 5.0 percentage points more negative than the corresponding interaction across the program choices. Omnibus measures consequently make opposition appear more categorical and ethnoracially divided than the underlying component decisions reveal. This comparison is descriptive rather than an identified causal effect of itemization, but it demonstrates that measurement level changes the political structure researchers observe.

The broader implication extends beyond foreign aid. Immigration, climate policy, health care, and national security all combine policies that differ in purpose, cost, and implementation. General questions about these domains may record opposition to the complete bundle while concealing support for particular components. They may also make political disagreement appear more fundamental than it is when citizens share a relative ordering but apply different acceptance thresholds. Choice structure consequently does more than affect the level of expressed support. It determines which coalitions are observable and what kind of disagreement researchers infer from public opinion.

This paper demonstrates that aggregation can conceal heterogeneity in both policies and

publics. Omnibus measures make conditional supporters observationally indistinguishable from categorical opponents, while average treatment effects can combine opposing responses to racial cues. Disaggregating policy components and respondent reactions reveals both the coalition hidden by the omnibus measure and the racial divisions hidden by the average.

References

- [1] Baker, Andrew. 2015. “Race, Paternalism, and Foreign Aid: Evidence from U.S. Public Opinion.” *American Political Science Review* 109 (1): 93–109. <https://doi.org/10.1017/S0003055414000549>.
- [2] Beauvais, Edana, and Adam M. Enders. 2026. “Toward a Theory and Measure of Racial Benevolence.” *Political Behavior*. <https://doi.org/10.1007/s11109-025-10117-4>.
- [3] Bizumic, Boris, Conal Monaghan, and Daniel Priest. 2021. “The Return of Ethnocentrism.” *Political Psychology* 42 (S1): 29–73. <https://doi.org/10.1111/pops.12710>.
- [4] Chicago Council on Global Affairs. 1974–2024. *Chicago Council Surveys of American Public Opinion and U.S. Foreign Policy*. Chicago: Chicago Council on Global Affairs.
- [5] Chong, Dennis, and James N. Druckman. 2007. “Framing Theory.” *Annual Review of Political Science* 10: 103–126. <https://doi.org/10.1146/annurev.polisci.10.072805.103054>.
- [6] Doherty, David, Amanda Clare Bryan, Dina Hanania, and Matthew Pajor. 2020. “The Public’s Foreign Aid Priorities: Evidence from a Conjoint Experiment.” *American Politics Research* 48 (5): 635–648. <https://doi.org/10.1177/1532673X20939925>.
- [7] Dworkin, Gerald. 1988. *The Theory and Practice of Autonomy*. Cambridge: Cambridge University Press.
- [8] Genovese, Eugene D., and Elizabeth Fox-Genovese. 2011. *Fatal Self-Deception: Slaveholding Paternalism in the Old South*. New York: Cambridge University Press.
- [9] Gilens, Martin. 2001. “Political Ignorance and Collective Policy Preferences.” *American Political Science Review* 95 (2): 379–396. <https://doi.org/10.1017/S0003055401002222>.

- [10] Glaser, James M. 2002. “White Voters, Black Schools: Structuring Racial Choices with a Checklist Ballot.” *American Journal of Political Science* 46 (1): 35–46. <https://doi.org/10.2307/3088412>.
- [11] Goetz, Rebecca Anne. 2012. *The Baptism of Early Virginia: How Christianity Created Race*. Baltimore: Johns Hopkins University Press.
- [12] Green-Riley, Naima, Joshua D. Kertzer, Chryl N. Laird, and Julian J. Wamble. 2026. “Black and White Beyond the Water’s Edge: Racial Gaps in Foreign Policy Preferences.” Forthcoming in *American Journal of Political Science*.
- [13] Guttman, Louis. 1950. “The Basis for Scalogram Analysis.” In *Measurement and Prediction*, eds. Samuel A. Stouffer et al. Princeton: Princeton University Press, 60–90.
- [14] Hainmueller, Jens. 2012. “Entropy Balancing for Causal Effects: A Multivariate Reweighting Method to Produce Balanced Samples in Observational Studies.” *Political Analysis* 20(1): 25–46.
- [15] Hainmueller, Jens, Jonathan Mummolo, and Yiqing Xu. 2019. “How Much Should We Trust Estimates from Multiplicative Interaction Models? Simple Tools to Improve Empirical Practice.” *Political Analysis* 27 (2): 163–192. <https://doi.org/10.1017/pan.2018.46>.
- [16] Hurst, Reuben, Taylor Tidwell, and Darren Hawkins. 2017. “Down the Rathole? Public Support for US Foreign Aid.” *International Studies Quarterly* 61 (2): 442–454. <https://doi.org/10.1093/isq/sqx019>.
- [17] Imai, Kosuke, Luke Keele, and Dustin Tingley. 2010. “A General Approach to Causal Mediation Analysis.” *Psychological Methods* 15(4): 309–334.
- [18] Inouye, Rikio. 2026. “More Equal than Others: Race and Religion Alter Support to Countries in Conflict.” APSA Preprints.

- [19] Jackman, Mary R. 1994. *The Velvet Glove: Paternalism and Conflict in Gender, Class, and Race Relations*. Berkeley: University of California Press.
- [20] Jacoby, William G. 2000. "Issue Framing and Public Opinion on Government Spending." *American Journal of Political Science* 44 (4): 750–767. <https://doi.org/10.2307/2669279>.
- [21] Jennings, Willie James. 2010. *The Christian Imagination: Theology and the Origins of Race*. New Haven, CT: Yale University Press.
- [22] Jordan, Winthrop D. 1974. *The White Man's Burden: Historical Origins of Racism in the United States*. New York: Oxford University Press.
- [23] Kates, Jennifer, Josh Michaud, Ashley Kirzinger, Bryan Wu, and Mollyann Brodie. 2019. *Where Does Public Opinion Stand on the U.S. Role in Global Health?* San Francisco: Kaiser Family Foundation. <https://www.kff.org/global-health-policy/poll-finding/where-does-public-opinion-stand-u-s-role-global-health/>.
- [24] Kinder, Donald R., and Cindy D. Kam. 2009. *Us Against Them: Ethnocentric Foundations of American Opinion*. Chicago: University of Chicago Press.
- [25] Kinder, Donald R., and David O. Sears. 1981. "Prejudice and Politics: Symbolic Racism versus Racial Threats to the Good Life." *Journal of Personality and Social Psychology* 40 (3): 414–431. <https://doi.org/10.1037/0022-3514.40.3.414>.
- [26] Milner, Helen V., and Dustin Tingley. 2013. "The Choice for Multilateralism: Foreign Aid and American Foreign Policy." *The Review of International Organizations* 8 (3): 313–341. <https://doi.org/10.1007/s11558-012-9153-x>.
- [27] Milner, Helen V., and Dustin Tingley. 2015. *Sailing the Water's Edge: The Domestic Politics of American Foreign Policy*. Princeton, NJ: Princeton University Press.

- [28] Moore, Steven. 2020. “The Road to Hell: How Race Paternalism Shapes Political Behavior.” Ph.D. dissertation, University of Michigan.
- [29] Mokken, Robert J. 1971. *A Theory and Procedure of Scale Analysis*. The Hague: Mouton.
- [30] Mutz, Diana C., Edward D. Mansfield, and Eunji Kim. 2021. “The Racialization of International Trade.” *Political Psychology* 42(4): 555–573. <https://doi.org/10.1111/pops.12714>.
- [31] Parajon, Eric. 2025. “The Effect of Racial Resentment and Out-Group Cues on Support for Climate Policy.” *International Studies Quarterly* 69 (4): sqaf087. <https://doi.org/10.1093/isq/sqaf087>.
- [32] Paxton, Pamela, and Stephen Knack. 2012. “Individual and Country-Level Factors Affecting Support for Foreign Aid.” *International Political Science Review* 33 (2): 171–192. <https://doi.org/10.1177/0192512111406095>.
- [33] Prather, Lauren. 2024. “Ideology at the Water’s Edge: Explaining Variation in Public Support for Foreign Aid.” *World Development* 176: 106472. <https://doi.org/10.1016/j.worlddev.2023.106472>.
- [34] Riker, William H. 1986. *The Art of Political Manipulation*. New Haven, CT: Yale University Press.
- [35] Scotto, Thomas J., Jason Reifler, David Hudson, and Jennifer vanHeerde-Hudson. 2017. “We Spend How Much? Misperceptions, Innumeracy, and Support for the Foreign Aid in the United States and Great Britain.” *Journal of Experimental Political Science* 4 (2): 119–128. <https://doi.org/10.1017/XPS.2017.6>.

- [36] Sheagley, Geoffrey, and Scott Clifford. 2025. “No Evidence That Measuring Moderators Alters Treatment Effects.” *American Journal of Political Science* 69(1): 49–63. <https://doi.org/10.1111/ajps.12814>.
- [37] Song, Lucy. 2026. “Does White identity stop at the water’s edge? White solidarity and racialized perceptions of foreign adversaries in the United States.” *Political Psychology* 47, (4): e70180. <https://doi.org/10.1111/pops.70180>
- [38] van Schuur, Wijbrandt H. 2003. “Mokken Scale Analysis: Between the Guttman Scale and Parametric Item Response Theory.” *Political Analysis* 11 (2): 139–163. <https://doi.org/10.1093/pan/mpg002>.
- [39] Westwood, Sean J. 2025. “The Potential Existential Threat of Large Language Models to Online Survey Research.” *Proceedings of the National Academy of Sciences* 122 (47): e2518075122.
- [40] Westwood, Sean J., and Samuel Frederick. 2026. “Reply to Van der Stigchel et al.: Empirical Evidence that AI Survey Contamination Is Real and Substantial.” *Proceedings of the National Academy of Sciences* 123 (8): e2537420123.
- [41] Williamson, Vanessa. 2019. “Public Ignorance or Elitist Jargon? Reconsidering Americans’ Overestimates of Government Waste and Foreign Aid.” *American Politics Research* 47(1): 152–173. <https://doi.org/10.1177/1532673X18759645>.

Supplementary Materials for Switching Sides: Choice Structure, Race, and American Support for Foreign Aid

September 28, 2026

1	Guide to the Supplement and Analysis Status	SI-3
2	Research Design, Sample, and Measurement	SI-4
3	Omnibus and Itemized Preferences (Main Text Section 7.1–7.2)	SI-8
3.1	Preregistered Comparisons and Program Allocations	SI-8
3.2	Itemized-Only-Arm Validation	SI-11
3.3	Common Ordering with Heterogeneous Cutpoints	SI-13
3.4	Covariate and Population-Composition Robustness	SI-14
4	Reported Program Evaluations and Delivery Preferences	SI-17
4.1	Program Evaluations by Ideology	SI-21
5	Country-Race Cue and Ethnocentric Moderation (Main Text Section 7.3)	SI-23
5.1	Average Effects of the Country-Race Cue	SI-23
5.2	Preregistered Moderation Tests	SI-25
5.3	Joint-Moderator Robustness	SI-27
5.4	Program-Level Moderation, Thresholds, and Ordering	SI-28
5.5	Linear-Interaction Diagnostics	SI-31
6	Country-Race Cue and Subsequent Program Evaluations (Main Text Section 7.4)	SI-33

7 Preregistered Pathway Diagnostics	SI-35
7.1 Sensitivity of the Pathway Estimates	SI-39
8 Additional Descriptive Analyses	SI-40
8.1 Correlates of Conditional Support	SI-40
8.2 Geographic References in Open-Ended Responses	SI-42
8.3 Perceived Region and Religion of the Recipient Country	SI-42
9 Survey Instrument	SI-44
9.1 Pre-Treatment Demographics	SI-44
9.2 Political Predispositions	SI-44
9.3 Treatment	SI-47
9.4 Manipulation Check	SI-47
9.5 Primary Outcomes	SI-47
9.6 Delivery Mode	SI-48
9.7 Open-Ended Item	SI-48
9.8 Post-Treatment Mediators	SI-49
9.9 Leader and Government Perceptions	SI-50
9.10 Mechanism Batteries	SI-51
9.11 Inference Checks	SI-51

1 Guide to the Supplement and Analysis Status

The supplement follows the order of the empirical results in the main text. The table below distinguishes preregistered tests from exploratory extensions, descriptive analyses, and non-preregistered robustness checks. A preregistered survey experiment does not imply that every analysis was preregistered; the status of each analysis is stated again at the beginning of the corresponding subsection.

Table A1: Guide to analysis status and supplementary evidence

Analysis family	Status	Location
Omnibus versus itemized support, program hierarchy, and allocations (H1–H3)	Preregistered tests	Main text Section 7.1; Supplement Section 3.1
Common cumulative ordering and heterogeneous cutpoints	Exploratory	Main text Section 7.2; Supplement Section 3.3
Program evaluations, delivery mode, and preregistered partisan predictions	Descriptive and preregistered components, identified within Section 4	Main text Section 7.2; Supplement Section 4
Average effects of recipient race	Exploratory; treatment is randomized	Main text Section 7.3.1; Supplement Section 5.1
Moderation of cash support and switching by ethnocentrism and racial resentment (H4–H5)	Preregistered tests	Main text Section 7.3.2; Supplement Section 5.2
Omnibus and pooled moderation, cross-outcome contrasts, joint moderators, rank stability, and funding thresholds	Exploratory	Main text Section 7.3.3; Supplement Section 5.3–5.4
Recipient race and subsequent program evaluations	Exploratory	Main text Section 7.4; Supplement Section 6
Paternalism, agency, and living-standards pathways (H6)	Preregistered pathway diagnostics; not causally identified mediation	Main text Section 7.5; Supplement Section 7
Population weighting and functional-form checks	Non-preregistered robustness analyses	Supplement Sections 3.4 and 5.5

Notes: “Exploratory” denotes analyses not specified in the preregistration. “Pathway diagnostics” reflects that the candidate mediators were measured after the funding outcomes. Full hypotheses and analysis specifications appear in the preregistration.

2 Research Design, Sample, and Measurement

Analysis status: Design information and measurement documentation; scale construction and exclusions follow the preregistration.

Table A2: Sample characteristics

Characteristic	All	Recipient country	
		White-majority	Black-majority
Age (years)	46.78 (14.51)	45.90 (14.44)	47.64 (14.53)
Female	60.9%	62.2%	59.7%
Bachelor's degree or higher	51.2%	50.6%	51.7%
Income bracket (1–14)	7.84 (3.13)	7.81 (3.12)	7.86 (3.14)
Evangelical	26.7%	25.9%	27.4%
Ideology (1 = liberal, 7 = conservative)	3.89 (1.88)	3.90 (1.88)	3.89 (1.87)
Democrat	35.1%	34.5%	35.6%
Republican	36.6%	37.4%	35.9%
Independent or other	28.3%	28.1%	28.5%
Internationalism (0–1)	0.59 (0.29)	0.59 (0.29)	0.58 (0.30)
Ethnocentrism (1–5)	2.04 (0.76)	2.03 (0.76)	2.05 (0.76)
Racial resentment (1–5)	2.86 (1.34)	2.87 (1.35)	2.85 (1.33)
Respondents	1,996	987	1,009

Notes: Cell entries are means with standard deviations in parentheses, except for binary characteristics, which are percentages. The sample is non-Hispanic white U.S. adults recruited on Prolific with quotas on self-reported ideology (35 percent liberal, 30 percent moderate, 35 percent conservative). Recipient race was randomly assigned; the condition columns are reported so that readers can see the composition of each arm. Ethnocentrism and racial resentment are measured before treatment. Data: Prolific, August 2026.

Table A3: Sample characteristics compared with White non-Hispanic adults

	Sample	ANES 2024	Difference
<i>Demographics (percent except age)</i>			
Male	38.9	49.0	-10.1
Mean age (years)	46.8	50.7	-3.9
Age 18–29	12.2	16.7	-4.5
Age 30–44	34.8	23.6	+11.2
Age 45–64	39.0	31.8	+7.2
Age 65+	14.0	27.9	-13.9
<i>Education (percent)</i>			
High school or less	15.4	34.2	-18.8
Some college or associate	33.4	29.3	+4.1
Bachelor’s degree	32.0	22.6	+9.4
Postgraduate degree	19.1	13.8	+5.3
<i>Ideology (percent)</i>			
Liberal	41.8	30.3	+11.5
Moderate	16.8	21.3	-4.4
Conservative	41.3	48.4	-7.1
<i>Party identification (percent)</i>			
Democrat	35.1	29.1	+6.0
Independent or other	28.3	32.8	-4.5
Republican	36.6	38.1	-1.5

Notes: The sample column describes the analysis sample ($N = 1,996$). The population column is White non-Hispanic adults in the ANES 2024 Time Series pre-election survey ($n = 3,946$), weighted with the full-sample pre-election weight; ANES ages are top-coded at 80. Ideology follows the paper’s three-category rule: placements 1–3 on the seven-point scale are liberal, 4 moderate, and 5–7 conservative, with ANES nonplacers assigned by the forced-choice follow-up. Party identification uses the initial party question; leaners are grouped with independents. The sample was recruited with quotas on three-category ideology (35 percent liberal, 30 percent moderate, 35 percent conservative), so ideological composition reflects design rather than self-selection alone. Data: Prolific, August 2026, and ANES 2024 Time Series.

Table A4: Attention and manipulation check pass rates

Check	All	Recipient country	
		White-majority	Black-majority
Attention check 1 (news frequency)	88.2% (2,026)	—	—
Attention check 2 (scale instruction)	99.9% (2,026)	—	—
Manipulation check (recipient race recall)	98.4% (1,996)	97.5% (987)	99.3% (1,009)

Notes: Entries are the share passing each check, with the number of respondents scored in parentheses. The two attention checks precede treatment; respondents who failed both were screened out before the vignette, so the rates shown are among respondents who reached treatment. The manipulation check is the post-treatment item, “What is the largest racial group in Glantonía?”; a pass names the group the vignette assigned. It is reported for the analysis sample after the preregistered keystroke exclusions. Data: Prolific, August 2026.

Table A5: Scale reliability

Scale	Role	Items	Timing	α
Ethnocentrism	Moderator (H4)	6	Pre-treatment	.77
Racial resentment	Moderator (H5)	4	Pre-treatment	.94
Paternalism (full)	Mediator (H6)	4	Post-treatment	.74
Benevolence sub-scale	Mediator (H6)	2	Post-treatment	.70
Control sub-scale	Mediator (H6)	2	Post-treatment	.69
Perceived agency (Baker)	Mediator (H6)	3	Post-treatment	.74
Perceived living standards (Baker)	Mediator (H6)	2	Post-treatment	.78

Notes: Cronbach's α , computed on the main arm ($N = 1,769$). For the two-item scales the inter-item correlations are .53 (benevolence), .53 (control), and .64 (living standards). The benevolence and control sub-scales correlate at $r = .46$, consistent with treating them as related but distinct dimensions of paternalism. Mediator scales are measured post-treatment, so their reliabilities are computed on the treated sample. Data: Prolific, August 2026.

3 Omnibus and Itemized Preferences (Main Text Section 7.1–7.2)

3.1 Preregistered Comparisons and Program Allocations

Analysis status: Preregistered tests of H1–H3.

Table A6: Between-arm test of question-order effects

Term	(1)	(2)
Saw omnibus first (order effect)	-0.023 (0.018) [0.19]	-0.024 (0.025) [0.34]
Black country	0.010 (0.012) [0.36]	0.008 (0.033) [0.82]
Order $\times$ Black country	—	0.003 (0.035) [0.93]
Intercept	0.519 (0.017) [0.00]	0.520 (0.024) [0.00]

Notes: OLS regressions of the share of the seven itemized programs supported on an indicator for the main arm (which answered the omnibus item before the checklist) and the Black-country treatment; column 2 adds their interaction. The itemized-only arm never saw the omnibus item. Cells report the estimate, HC2 robust standard error in parentheses, and p-value in brackets. N = 1,996 (1,769 main arm; 227 itemized-only). Data: Prolific, August 2026.

Table A7: Support for the aid package as a whole and for its component programs

Program	Share	95% CI
Medicine	86.7%	[0.851, 0.882]
Food	86.0%	[0.844, 0.876]
Transportation infrastructure	60.7%	[0.584, 0.630]
Democratic elections	51.7%	[0.493, 0.540]
Cash to the citizens	37.0%	[0.348, 0.393]
Military equipment	16.2%	[0.145, 0.179]
Cash to the government	12.9%	[0.113, 0.145]
Omnibus package	36.4%	—

Notes: Share of main-arm respondents supporting funding for each program on the itemized check-list; the final row is support for the identical \$50 million package on the single omnibus item. Confidence intervals from the normal approximation to the binomial. Main-arm N = 1,769. Data: Prolific, August 2026.

Table A8: Program choices of omnibus supporters, soft switchers, and switchers

Program	Omnibus supporters	Soft switchers	Switchers
Medicine	99.8%	98.7%	94.0%
Food	99.4%	97.7%	93.1%
Transportation infrastructure	89.9%	76.8%	41.3%
Democratic elections	71.9%	67.5%	38.9%
Cash to the citizens	61.8%	45.7%	18.5%
Cash to the government	26.2%	12.2%	3.4%
Military equipment	25.6%	18.0%	10.5%

Notes: Share of each group supporting funding for the program. Omnibus supporters (n = 644) back the package on the up-or-down vote; soft switchers (n = 311) answer neither and fund at least one program; switchers (n = 620) oppose the package and fund at least one program. Data: Prolific, August 2026.

3.2 Itemized-Only-Arm Validation

Analysis status: Exploratory; not preregistered.

The itemized-only arm provides an independent assessment of the program hierarchy among respondents who did not first evaluate the omnibus proposal. Medicine and food receive the greatest support, followed by transportation infrastructure and democratic elections, while military equipment and cash transfers receive substantially less support. Table A9 reports these patterns by country-race condition.

Table A9: Program support by recipient race in the itemized-only arm

Program	White-majority	Black-majority	Diff.	(SE)	<i>p</i>
Medicine	87.0	90.2	3.2	(4.2)	0.447
Food	87.0	88.4	1.4	(4.4)	0.743
Military equipment	20.9	9.8	−11.0	(4.7)	0.021
Democratic elections	48.7	56.2	7.6	(6.6)	0.256
Transportation infrastructure	65.2	71.4	6.2	(6.2)	0.317
Cash to the government	14.8	13.4	−1.4	(4.6)	0.765
Cash to the citizens	40.9	40.2	−0.7	(6.5)	0.916

Notes: Exploratory. Share of itemized-only respondents supporting funding for each program, by recipient-race condition (White-majority $n = 115$; Black-majority $n = 112$). This arm answered the program checklist without first seeing the omnibus item. Differences are from per-program OLS models with HC2 standard errors, unadjusted for multiple comparisons; no per-program difference survives a Benjamini–Hochberg correction across the seven programs. A joint test cannot reject that recipient race leaves all seven programs unchanged ($\chi^2(7) = 9.65$, $p = 0.21$). Data: Prolific, August 2026.

Table A10: Program support and allocations by respondent group

Program	Omnibus supporters		Soft switchers		Switchers	
	Support	Alloc. (%)	Support	Alloc. (%)	Support	Alloc. (%)
Medicine	99.8%	26.9	98.7%	29.4	94.0%	36.5
Food	99.4%	30.0	97.7%	31.5	93.1%	38.0
Transportation infrastructure	89.9%	14.9	76.8%	13.5	41.3%	9.0
Democratic elections	71.9%	10.5	67.5%	11.0	38.9%	7.5
Cash to the citizens	61.8%	10.2	45.7%	8.6	18.5%	5.0
Military equipment	25.6%	3.8	18.0%	3.5	10.5%	2.8
Cash to the government	26.2%	3.7	12.2%	2.4	3.4%	1.1

Notes: Appendix version of Figure 1. Support is the share of the group funding each program on the itemized checklist; Alloc. is the mean percentage of the \$50 million budget the group assigns to the program in the allocation task. Omnibus supporters ($n = 644$) back the package on the up-or-down vote; soft switchers ($n = 311$) answer neither and fund at least one program; switchers ($n = 620$) oppose the package and fund at least one program. Data: Prolific, August 2026.

3.3 Common Ordering with Heterogeneous Cutpoints

Analysis status: Exploratory; not preregistered.

If respondents share a common ordering over aid instruments and differ primarily in their acceptance thresholds, their choices should approximate a cumulative Guttman pattern (Guttman 1950). I evaluate this structure using the coefficient of reproducibility and Loevinger’s H , a measure of scalability from nonparametric Mokken scale analysis (Mokken 1971; van Schuur 2003).


Table A11: The seven program choices form a common ordering with heterogeneous cutpoints

<i>Panel A. Scalability of the choice sets</i>	
Choice sets exactly matching an ordering-plus-cutpoint pattern	64.2%
Expected under independent choices	35.9%
Choice sets within one deviation of such a pattern	95.4%
Choice sets within two deviations	99.8%
Coefficient of reproducibility	0.942
Loevinger’s H (full scale)	0.615
<i>Panel B. Cutpoint distribution, exact patterns only (programs funded)</i>	
0 programs	17.1%
1 program	1.7%
2 programs	19.8%
3 programs	10.3%
4 programs	22.1%
5 programs	21.4%
6 programs	2.2%
7 programs	5.5%

Notes: The modal ordering ranks the seven programs by overall support (medicine, food, transportation infrastructure, democratic elections, cash to the citizens, military equipment, cash to the government). A choice set matches an ordering-plus-cutpoint pattern if the respondent supports every program above some personal threshold in this ordering and rejects every program below it; deviations are the minimum number of choices that would have to change to produce such a pattern. The independence benchmark simulates 20,000 respondents choosing each program independently at the observed marginal rates. The coefficient of reproducibility is one minus total deviations over total choices; Loevinger’s H compares observed pairwise Guttman errors to those expected under independence, with $H > .50$ conventionally read as a strong cumulative scale. Panel B shows that respondents whose choices fit the scale exactly spread across the full range of cutpoints. Main arm, $N = 1,769$. Data: Prolific, August 2026.

3.4 Covariate and Population-Composition Robustness

Analysis status: Non-preregistered robustness analyses.



Notes: Appendix version of Figure 1. Each point is the predicted probability of supporting funding for the named program at the covariate means, from a per-program linear probability model controlling for ideology (mean-centered), Independent and Republican indicators, internationalism, self-monitoring, evangelical identification, age, gender, education, and income (HC2 standard errors). The dashed line is the unadjusted share supporting the omnibus package. Cash programs in crimson. Thick and thin whiskers denote 90% and 95% confidence intervals. Main-arm N = 1,769. Data: Prolific, August 2026.

Table A12: Support for the component programs, adjusted for covariates

Program	Adjusted share	95% CI
Medicine	86.7%	[0.852, 0.881]
Food	86.0%	[0.845, 0.875]
Transportation infrastructure	60.7%	[0.587, 0.628]
Democratic elections	51.7%	[0.496, 0.537]
Cash to the citizens	37.0%	[0.349, 0.391]
Military equipment	16.2%	[0.145, 0.178]
Cash to the government	12.9%	[0.114, 0.144]

Notes: Predicted probability of supporting each program at the covariate means, from per-program linear probability models controlling for ideology (mean-centered), Independent and Republican indicators, internationalism, self-monitoring, evangelical identification, age, gender, education, and income (HC2 standard errors). Main-arm N = 1,769. Data: Prolific, August 2026.

Table A13: Headline quantities reweighted to white adult demographic margins


Quantity	Unweighted	Balanced	Difference
<i>Headline quantities</i>			
Support the omnibus package	36.4	35.9	-0.5
Fund at least one program	89.0	89.0	+0.0
Switch (among omnibus opponents)	76.4	76.8	+0.3
<i>Program support (Figure 1)</i>			
Medicine	86.7	86.2	-0.4
Food	86.0	85.6	-0.4
Military equipment	16.2	16.4	+0.2
Democratic elections	51.7	51.7	+0.0
Transportation infrastructure	60.7	60.6	-0.1
Cash to the government	12.9	11.6	-1.3
Cash to the citizens	37.0	36.1	-0.9
<i>Program beliefs, averaged over programs (Figure A2)</i>			
Saves lives	59.4	58.8	-0.6
Morally right	57.5	57.2	-0.3
Easy to verify	45.9	45.8	-0.1
Reduces theft	33.5	33.5	+0.0
Increases U.S. influence	49.9	50.7	+0.8
Benefits U.S. economy	43.9	44.3	+0.5

Notes: Exploratory; not preregistered. Entries are percentages. “Balanced” applies entropy balancing weights (Hainmueller 2012) that set the sample’s sex, age, and education shares equal to national margins for non-Hispanic white adults age 18 and over. Sex and age margins are from the 2020–2024 American Community Survey five-year estimates (table B01001H); education margins are from the Current Population Survey (NCES Digest of Education Statistics 2023, Table 104.30). The constraints are matched to within 10^{-4} . Balancing uses demographic characteristics only, because the attitudinal measures have no external benchmark. The sample is quota-sampled on ideology to power the moderation tests and is not designed to reproduce population levels; this table reports how far the headline shares depend on demographic composition, not a population-representative estimate. Main arm, $N = 1,769$. Data: Prolific, August 2026.

4 Reported Program Evaluations and Delivery Preferences

Analysis status: Descriptive analyses plus preregistered predictions concerning delivery mode, military assistance, and partisan differences; the text and notes identify each test.

Figure A2: Beliefs about each aid program



Notes: Each panel shows respondents’ evaluations of one of the seven aid programs. Bars report the percentage of main-arm respondents who indicate that the program saves lives, is morally right, increases U.S. influence, is easy to verify, benefits the U.S. economy, or reduces theft. Main-arm $N = 1,769$. Data: Prolific, August 2026.

Table A14: Beliefs about each aid program (percent answering Yes)

Program	Theft	Verify	Economy	Influence	Lives	Moral
Medicine	47	73	56	53	98	96
Food	49	69	52	49	98	96
Military equipment	15	36	41	63	34	23
Democratic elections	50	43	52	64	48	59
Transportation infrastructure	31	61	53	50	64	67
Cash to the government	9	12	20	42	14	12
Cash to the citizens	34	27	33	28	60	50

Notes: Share of main-arm respondents answering Yes when asked whether the program reduces theft, is easy to verify, benefits the U.S. economy, increases U.S. influence, saves lives, and is morally right. Entries are percentages. Main-arm N = 1,769. Data: Prolific, August 2026.

Table A15: Item-level support and mechanism beliefs

Belief	(1) Baseline	(2) Military	(3) Partisan
Reduces theft	0.071*** (0.010)	0.060*** (0.010)	0.059*** (0.015)
Easy to verify	0.123*** (0.011)	0.131*** (0.011)	0.115*** (0.015)
Morally right	0.476*** (0.014)	0.446*** (0.015)	0.507*** (0.025)
Saves lives	0.173*** (0.013)	0.169*** (0.012)	0.177*** (0.023)
Increases U.S. influence	0.033*** (0.009)	0.061*** (0.010)	0.017 (0.014)
Military program	—	-0.098*** (0.011)	—
Influence $\times$ Military	—	-0.079*** (0.016)	—
Republican	—	—	-0.064*** (0.015)
Theft $\times$ Rep.	—	—	0.012 (0.024)
Verify $\times$ Rep.	—	—	0.000 (0.025)
Moral $\times$ Rep.	—	—	-0.088*** (0.033)
Lives $\times$ Rep.	—	—	-0.011 (0.030)
Influence $\times$ Rep.	—	—	0.030 (0.021)

Notes: Linear probability models of support for each itemized program (one row per respondent $\times$ program; standard errors clustered by respondent in parentheses). Column 1 is the baseline model on the full main arm ($N = 1,769 \times 7$). Column 2 adds a military-program indicator and its interaction with the influence belief. Column 3 interacts every belief with Republican identification, restricted to partisans (Democratic baseline). * $p < .10$, ** $p < .05$, *** $p < .01$. Data: Prolific, August 2026.

Table A16: Mean rank of each delivery mode, by program


Program	Training	Cash	Supplies	p (cash vs. best alt.)
Medicine	2.04	2.58	1.39	< .001
Food	2.38	2.26	1.36	< .001
Military equipment	1.39	2.71	1.90	< .001
Transportation infrastructure	1.95	2.30	1.75	< .001
Democratic elections	1.31	2.54	2.15	< .001

Notes: Mean respondent ranking of the three delivery modes for each of the five non-cash programs (1 = best, 3 = worst). The final column is the p-value from a paired t-test of the cash rank against each respondent's better-ranked alternative mode. Main-arm N = 1,769. Data: Prolific, August 2026.

4.1 Program Evaluations by Ideology

Analysis status: Preregistered exploratory predictions.

Figure A3: Beliefs about each aid program, by respondent ideology



Notes: Each panel shows one of the seven aid programs. Points give the share of liberal (blue circles, $n = 711$) and conservative (crimson triangles, $n = 757$) main-arm respondents answering Yes to each mechanism question; whiskers denote 95% confidence intervals. Moderates omitted for legibility. Data: Prolific, August 2026.

Table A17: Beliefs about each aid program, by respondent ideology (percent Yes, Liberal/Conservative)

Program	Theft	Verify	Economy	Influence	Lives	Moral
Medicine	56/41	79/68	63/49	58/50	99/97	99/94
Food	55/43	76/64	58/46	53/46	100/96	99/94
Military equipment	13/18	36/38	39/43	67/61	26/43	15/29
Democratic elections	65/39	54/33	65/42	74/57	66/37	77/45
Transportation infrastructure	37/27	72/54	60/45	56/46	76/55	78/56
Cash to the government	9/9	12/12	24/16	51/34	18/11	14/10
Cash to the citizens	44/26	33/21	39/25	33/21	77/44	66/36


Notes: Each cell gives the share of liberal ($n = 711$) and conservative ($n = 757$) main-arm respondents answering Yes, as Liberal/Conservative percentages. Moderates omitted. Data: Prolific, August 2026.

5 Country-Race Cue and Ethnocentric Moderation (Main Text Section 7.3)

5.1 Average Effects of the Country-Race Cue

Analysis status: Exploratory; not preregistered. Treatment assignment is randomized.

Figure A4: Support for the Omnibus Proposal and Individual Programs by Country-Race Condition



Notes: Each point is the share of main-arm respondents supporting the omnibus proposal (top row) or funding the named program on the itemized checklist, by country-race condition (white-country: crimson triangles, $n = 872$; Black-country: black circles, $n = 897$). Thick and thin whiskers denote 90% and 95% confidence intervals. The omnibus contrast is the exploratory average effect reported in the text; the program-level contrasts are small and mostly indistinguishable from zero (Appendix Table A18). Data: Prolific, August 2026.

Table A18 reports the average effect of assignment to the Black-country condition on each outcome together with two one-sided equivalence tests. The omnibus effect is large, whereas several aggregate itemized outcomes can be bounded within substantively small equivalence margins.

Table A18: No Controls: Black–White Difference, with equivalence tests

Outcome	B-W Difference	SE	p	Equivalence p	
				$\pm 3pp$	$\pm 5pp$
Omnibus package	0.114	0.023	< .001	1.00	1.00
Fund ≥ 1 program	0.001	0.015	.955	.025	< .001
Item-support share	0.011	0.012	.377	.060	< .001
Humanitarian support	0.019	0.016	.236	.233	.022
Cash support	0.010	0.016	.508	.108	.006
Switching (opponents)	−0.032	0.030	.292	.521	.269
Medicine	0.017	0.016	.284	.217	.022
Food	0.020	0.017	.231	.268	.034
Military equipment	−0.038	0.018	.028	.686	.256
Democratic elections	0.017	0.024	.473	.293	.083
Transportation infrastructure	0.039	0.023	.090	.657	.324
Cash to the government	0.033	0.016	.041	.564	.136
Cash to the citizens	−0.012	0.023	.614	.212	.047

Notes: Each row reports the difference in the named outcome between the Black-majority and White-majority conditions from an unadjusted OLS model with HC2 standard errors; outcomes are scaled 0–1. Equivalence tests use two one-sided tests (TOST) with bounds of ± 3 and ± 5 percentage points. Values below .05 reject effects outside the stated equivalence bound. The omnibus effect is essentially unchanged with the full covariate set (0.113, SE 0.021) and is moderated by ethnocentrism (Black-majority $\times$ ethnocentrism = −0.093, SE 0.021, $p < .001$). The switching estimate is a descriptive Black–White contrast among observed omnibus opponents rather than a treatment effect because omnibus opposition is post-treatment. The opponent pools differ across conditions in both size (48.4 percent White-majority; 43.4 percent Black-majority) and ethnocentrism (Black-majority opponents score 0.18 points higher, $p = .001$). Switching is calculated among omnibus opponents ($n = 811$); all other rows use the main arm ($N = 1,769$). Data: Prolific, August 2026.

5.2 Preregistered Moderation Tests


Analysis status: Preregistered tests of H4 and H5. Switching is defined using a post-treatment omnibus response and is interpreted descriptively.

Table A19: Ethnocentrism and racial resentment as moderators of the recipient-race effect

Interaction	Outcome	Estimate	(SE)	<i>p</i>
Black country x Ethnocentrism	Omnibus support	-0.090	(0.020)	<0.001
Black country x Racial resentment	Omnibus support	-0.097	(0.020)	<0.001
Black country x Ethnocentrism	Switching (among opponents)	-0.055	(0.028)	0.051
Black country x Ethnocentrism	Cash-item support	-0.033	(0.014)	0.021
Black country x Racial resentment	Switching (among opponents)	-0.029	(0.033)	0.375
Black country x Racial resentment	Cash-item support	-0.018	(0.015)	0.209

Notes: Black-country x moderator interactions from separate OLS models; moderators standardized. The switching and cash-item rows are the pre-registered H5/H6 tests; the omnibus rows are exploratory. The omnibus outcome is coded 1 for respondents who support or somewhat support the package. All models control for ideology (mean-centered), Independent and Republican indicators, internationalism, self-monitoring, evangelical identification, age, gender, education, and income. HC2 robust standard errors in parentheses. Data: Prolific, August 2026.

Figure A5: Ethnoracial Attitudes as Moderators of the Country-Race Cue: Cash and Switching



Notes: Each point reports the interaction between assignment to the Black-country condition and the named moderator from a separate OLS model testing preregistered Hypotheses 4 and 5. Moderators are standardized, so the coefficients indicate how much the effect of the Black-country condition changes with a one-standard-deviation increase in the moderator. Models adjust for ideology, party identification, internationalism, self-monitoring, evangelical identification, age, gender, education, and income. Thick and thin whiskers denote 90% and 95% confidence intervals based on HC2 robust standard errors. Switching is measured among respondents who oppose the omnibus proposal and should therefore be interpreted descriptively rather than as a treatment effect. Data: Prolific, August 2026.

5.3 Joint-Moderator Robustness

Analysis status: Exploratory; not preregistered.

Because ethnocentrism and racial resentment are correlated ($r = .61$), Table A20 estimates their interactions with the country-race cue simultaneously. Conditional on racial resentment, the ethnocentrism interaction remains negative for cash-program support (-0.039 , $p = .029$) and pooled program support (-0.048 , $p = .002$), whereas the corresponding racial-resentment interactions are close to zero. The differences between the two interactions are not statistically significant for cash support ($p = .175$) and reach only marginal significance for pooled program support ($p = .058$). These results do not establish that ethnocentrism is the stronger moderator. They indicate that ethnocentrism captures variation in program-level responses to the country-race cue that is not absorbed by racial resentment.

The omnibus outcome does not exhibit the same distinction. When both moderators enter simultaneously, their interactions with the country-race cue are negative and statistically indistinguishable from one another ($p = .858$). Evidence differentiating the two moderators is therefore concentrated in respondents' program choices rather than their omnibus evaluations.

Table A20: Ethnocentrism and Racial Resentment as Joint Moderators of the Country-Race Cue

Outcome	$T \times$ Ethnocentrism		$T \times$ Resentment		Difference		p
	Est.	(SE)	Est.	(SE)	Est.	(SE)	
Omnibus support	-0.055	(0.025)	-0.063	(0.026)	0.008	(0.046)	0.858
Cash-item support	-0.039	(0.018)	0.005	(0.018)	-0.044	(0.032)	0.175
Pooled program support	-0.048	(0.015)	0.003	(0.014)	-0.051	(0.027)	0.058

Notes: Exploratory; not preregistered. Each row is one OLS model containing assignment to the Black-country condition, standardized ethnocentrism, standardized racial resentment, and both treatment-by-moderator interactions ($r = .61$ between the moderators). The final columns test the equality of the two interactions. The omnibus and cash models include the full covariate set (ideology, mean-centered, Independent and Republican indicators, internationalism, self-monitoring, evangelical identification, age, gender, education, and income) with HC2 standard errors. The pooled model stacks the seven program choices with program fixed effects and standard errors clustered by respondent. The omnibus outcome is coded 1 for respondents who support or somewhat support the proposal. Main arm, $N = 1,769$. Data: Prolific, August 2026.

5.4 Program-Level Moderation, Thresholds, and Ordering

Analysis status: Exploratory; not preregistered.

The following analyses evaluate whether ethnocentric moderation is specific to particular programs or instead moves respondents’ funding thresholds across a largely shared hierarchy. I first estimate the country-race-by-ethnocentrism interaction separately for each program and in a pooled model. I then test whether the program-specific interactions differ, whether the interaction predicts the number of programs funded, and whether the observed interaction is stronger for the omnibus evaluation than for the program choices.

Table A21: Program-Specific Moderation of the Country-Race Cue by Ethnocentrism

Program	Cue $\times$ Ethnocentrism	SE	p
Medicine	−0.047	0.018	.010
Food	−0.041	0.019	.026
Military equipment	−0.057	0.018	.001
Democratic elections	−0.043	0.022	.051
Transportation infrastructure	−0.050	0.021	.016
Cash to the government	−0.038	0.015	.013
Cash to the citizens	−0.022	0.021	.312
Pooled (all seven programs)	−0.043	0.012	< .001

Notes: Exploratory; not preregistered. Each program row reports a separate OLS model of program support on assignment to the Black-country condition, standardized ethnocentrism, and their interaction. The interaction indicates how much the effect of the Black-country condition changes with a one-standard-deviation increase in ethnocentrism. Standard errors are HC2. The pooled row stacks the seven programs (one row per respondent $\times$ program) with program fixed effects and standard errors clustered by respondent. Program-specific p -values are not adjusted for multiple comparisons. Main-arm $N = 1,769$. Data: Prolific, August 2026.

Table A22: Program Rankings across Country-Race and Ethnocentrism Profiles

Program	Lower ethnocentrism (−1 SD)		Higher ethnocentrism (+1 SD)	
	White-country	Black-country	White-country	Black-country
Medicine	89.2 (2)	95.6 (1)	83.3 (1)	83.5 (1)
Food	89.4 (1)	95.4 (2)	81.9 (2)	82.6 (2)
Transportation infrastructure	63.1 (3)	72.3 (3)	55.1 (3)	53.9 (3)
Democratic elections	57.3 (4)	61.7 (4)	46.2 (4)	42.2 (4)
Cash to the citizens	37.7 (5)	40.4 (5)	36.6 (5)	32.8 (5)
Military equipment	13.0 (6)	14.4 (7)	22.6 (6)	14.2 (6)
Cash to the government	9.7 (7)	16.5 (6)	13.1 (7)	12.4 (7)

Notes: Entries are predicted support in percent, with each program’s within-profile rank in parentheses (1 = most supported). Predictions are average predicted probabilities from the logistic models underlying Figure 2. Each model contains assignment to the Black-country condition, continuous ethnocentrism, their interaction, and the full covariate set; predictions are averaged over the observed covariate distribution. The six pairwise Spearman rank correlations are .93 (White-low, Black-low), .96 (White-low, White-high), .96 (White-low, Black-high), .96 (Black-low, White-high), .96 (Black-low, Black-high), and 1.00 (White-high, Black-high). Programs are ordered by overall support. Main arm, $N = 1,769$. Data: Prolific, August 2026.

Table A23: The Country-Race Cue and Ethnocentrism Move Funding Thresholds, Not Program Orderings

Quantity	Estimate	(SE)	<i>p</i>
Equality of the seven program interactions	$\chi^2(6) = 3.13$		0.793
Equivalence of all seven deviations within ± 0.05	Max. deviation = 0.021		0.041
Country-race cue $\times$ ethnocentrism, programs funded (0–7)	–0.299	(0.086)	< .001
with covariates	–0.279	(0.078)	< .001
Omnibus – programs difference in interaction	–0.050	(0.019)	0.007


Notes: Exploratory; not preregistered. The first row reports a Wald test of equality across the seven program-specific interactions. The second reports an intersection-union equivalence test using ± 5 -percentage-point bounds; the largest estimated deviation from the mean interaction is 2.1 points, and the reported *p*-value is the largest of the seven TOST *p*-values. The next two rows report OLS models of the number of programs funded, without and with the full covariate set. The final row reports the three-way interaction comparing the omnibus and program-level outcomes in a stacked model with standard errors clustered by respondent. This comparison is descriptive rather than a causal effect of itemization. Ethnocentrism is standardized. Main arm, $N = 1,769$. Data: Prolific, August 2026.

5.5 Linear-Interaction Diagnostics

Analysis status: Non-preregistered robustness analysis of the functional-form assumption.


Following Hainmueller, Mummolo, and Xu (2019), the country-race cue $\times$ ethnocentrism interactions are re-estimated without imposing a linear interaction. Figure A6 compares tercile-specific estimates with the conditional effects implied by the linear model. Figure A7 presents kernel estimates with cross-validated bandwidths. Both analyses are broadly consistent with an increasingly negative treatment effect as ethnocentrism rises, although uncertainty increases in the upper tail.

Figure A6: Effect of the Black-Country Condition by Ethnocentrism Tercile



Notes: Each point is the effect of assignment to the Black-country condition estimated separately within a tercile of ethnocentrism, plotted at the tercile median (whiskers denote 95% confidence intervals, HC2 standard errors); the dashed line is the conditional effect implied by the linear interaction model. Cash support uses the main arm ($N = 1,769$); switching uses omnibus opponents ($n = 811$). Data: Prolific, August 2026.

Figure A7: Effect of the Black-Country Condition across the Ethnocentrism Range



Notes: Each panel shows the conditional effect of assignment to the Black-country condition across the ethnocentrism range from the kernel estimator of Hainmueller, Mummolo & Xu (2019) (*interflex*), with cross-validated bandwidths (2.38 for cash support, 1.26 for switching) and 200 bootstrap replications; shaded bands denote pointwise 95% confidence intervals. Models include the full covariate set. The cash-support estimate declines across the observed range. The switching estimate is nonmonotonic near the lower bound but becomes increasingly negative as ethnocentrism rises above its mean. Cash support uses the main arm ($N = 1,769$); switching uses omnibus opponents ($n = 811$). Data: Prolific, August 2026.

6 Country-Race Cue and Subsequent Program Evaluations (Main Text Section 7.4)

Analysis status: Exploratory; not preregistered. Program evaluations were measured after the funding choices and are not interpreted as mediators.

Table A24: The recipient-race cue and program beliefs

Outcome	Average effect		$\times$ Ethnocentrism		At -1 SD		At $+1$ SD	
	Est.	p	Est.	p/q	Est.	p	Est.	p
Affirmative-evaluation index	+0.1	.940	−3.8	<.001	+4.0	<.001	−3.6	.012
<i>Component beliefs</i>								
Easy to verify	−0.6	.660	−4.3	.003	+4.0	.023	−4.7	.014
Increases U.S. influence	−1.2	.440	−4.1	.012	+3.0	.173	−5.2	.022
Reduces theft	+2.0	.128	−4.0	.004	+6.2	<.001	−1.9	.319
Morally right	−0.3	.752	−3.8	.003	+3.7	.006	−3.9	.021
Benefits U.S. economy	+0.8	.607	−3.6	.024	+4.6	.040	−2.6	.254
Saves lives	−0.3	.754	−2.9	.007	+2.8	.030	−3.0	.046

Notes: Entries are percentage points. The index is the mean of the six belief items; it has a Cronbach’s α of 0.75 (0.69 after removing program means; 0.55 to 0.77 within programs). Each outcome is estimated by stacking the seven programs (one row per respondent $\times$ program) with program fixed effects and standard errors clustered by respondent. “Average effect” is the effect of assignment to the Black-majority condition, estimated without the interaction; the conditional estimates come from the interaction model, whose treatment coefficient is the effect at mean ethnocentrism rather than the sample average. For the index, the interaction column reports a p -value; for the six component beliefs it reports Benjamini–Hochberg q -values adjusting for the six-test family. A joint Wald test that all six interactions equal zero rejects ($\chi^2(6) = 17.5$, $p = 0.008$); a test that they are equal to one another does not ($\chi^2(5) = 1.9$, $p = 0.86$): the six interaction estimates do not differ detectably from one another, supporting the common index as a summary of the pattern. Failure to reject equality does not establish that the six effects are identical. The joint model absorbs all 42 belief-by-program cells and allows belief-specific treatment and moderator terms. The conditional estimates one standard deviation below and above mean ethnocentrism are pointwise and unadjusted for multiple comparisons. The belief items are asked after the funding choices, so these are effects on reported beliefs rather than mediating pathways. Main-arm $N = 1,769$; 12,383 respondent-program rows per model. Data: Prolific, August 2026.

Table A25: Effect of the recipient-race cue on each program belief

Program	Saves lives	Morally right	Easy to verify	Reduces theft	Increases influence	Benefits economy
Medicine	-0.3 (0.7)	-0.6 (1.0)	+1.8 (2.1)	+4.4 (2.4)	-0.5 (2.4)	+1.9 (2.4)
Food	+0.3 (0.7)	+0.2 (0.9)	-0.8 (2.2)	+5.2* (2.4)	+1.5 (2.4)	+3.2 (2.4)
Military equipment	-3.1 (2.2)	-5.2* (2.0)	-4.0 (2.3)	-1.6 (1.7)	-6.6* (2.3)	-1.0 (2.3)
Democratic elections	+1.1 (2.4)	+3.1 (2.3)	+2.9 (2.4)	+2.2 (2.4)	-0.1 (2.3)	+0.2 (2.4)
Transportation infrastructure	+0.4 (2.3)	+1.8 (2.2)	+0.5 (2.3)	+4.7* (2.2)	+1.6 (2.4)	+3.7 (2.4)
Cash to the government	+2.1 (1.7)	+1.1 (1.6)	+0.1 (1.6)	+1.2 (1.4)	-3.8 (2.3)	-1.9 (1.9)
Cash to the citizens	-2.7 (2.3)	-2.9 (2.4)	-4.5* (2.1)	-2.1 (2.3)	-0.7 (2.1)	-0.2 (2.2)

Notes: Exploratory; not preregistered. Each entry is the effect of assignment to the Black-majority condition on the share answering Yes to that belief about that program, in percentage points, with heteroskedasticity-robust standard errors in parentheses. Each estimate comes from a separate regression on the single item, so there is one observation per respondent and no clustering is required. Of the forty-two estimates, five reach $p < .05$ uncorrected (*), and none survives a Benjamini–Hochberg correction across the forty-two tests (smallest $q = 0.16$); at $\alpha = .05$, roughly two would be expected by chance alone. The evidence for treatment effects on beliefs is in the interaction with ethnocentrism reported in Table A24, not in average effects on individual items. Main arm, $N = 1,769$. Data: Prolific, August 2026.

7 Preregistered Pathway Diagnostics


Analysis status: Preregistered tests of H6. Because candidate mediators were measured after the outcomes, the estimates are treated as pathway diagnostics rather than identified causal mediation effects.

Table A26: Estimated mediation of the recipient-race effect, by mediator and outcome

Mediator	Outcome	ACME	95% CI
Paternalism (full)	Omnibus support	0.0049	[-0.003, 0.014]
Paternalism (full)	Cash-item support	0.0039	[-0.002, 0.010]
Paternalism (full)	Humanitarian-item support	0.0048	[-0.003, 0.013]
Benevolence sub-scale	Omnibus support	0.0097	[-0.000, 0.020]
Benevolence sub-scale	Cash-item support	0.0075	[-0.000, 0.015]
Benevolence sub-scale	Humanitarian-item support	0.0085	[-0.000, 0.018]
Control sub-scale	Omnibus support	0.0003	[-0.003, 0.004]
Control sub-scale	Cash-item support	0.0003	[-0.003, 0.004]
Control sub-scale	Humanitarian-item support	0.0005	[-0.004, 0.005]
Perceived agency (Baker)	Omnibus support	0.0056	[-0.000, 0.012]
Perceived agency (Baker)	Cash-item support	0.0062	[-0.000, 0.013]
Perceived agency (Baker)	Humanitarian-item support	0.0063	[-0.000, 0.013]
Perceived living standards (Baker)	Omnibus support	0.0045	[0.002, 0.008]
Perceived living standards (Baker)	Cash-item support	0.0064	[0.003, 0.011]
Perceived living standards (Baker)	Humanitarian-item support	0.0071	[0.003, 0.012]

Notes: Average causal mediation effects of the Black-country treatment, from the pre-registered paired-equation models estimated by simulation. Outcomes are scaled 0–1. Both equations control for ideology (mean-centered), Independent and Republican indicators, internationalism, self-monitoring, evangelical identification, age, gender, education, and income. Data: Prolific, August 2026.

Figure A8: Model-Based Indirect-Effect Estimates by Candidate Pathway and Outcome




Notes: Each point reports the estimated average causal mediation effect (ACME) of assignment to the Black-majority condition through the named post-treatment mediator. Estimates and confidence intervals are expressed in percentage points. ACMEs come from the preregistered paired-equation models and are estimated using quasi-Bayesian simulation with 1,000 draws (Imai et al. 2010). Bars denote 95% confidence intervals. Both equations adjust for ideology, party identification, internationalism, self-monitoring, evangelical identification, age, gender, education, and income. Identification relies on sequential ignorability. The mediators were measured after the outcomes; the estimates should therefore be interpreted as evidence consistent with the proposed pathways rather than definitive evidence of causal mediation. Data: Prolific, August 2026.

Table A27: Joint mediator model: switching on all candidate mediators

Mediator	Estimate	(SE)	<i>p</i>
Paternalism	0.104	(0.019)	0.000
Perceived agency (Baker)	-0.058	(0.017)	0.001
Perceived living standards (Baker)	-0.085	(0.074)	0.251

Notes: OLS regression of switching (support for at least one itemized program among omnibus opponents, $n = 811$) on all candidate mediators entered simultaneously, with the full covariate set (ideology, partisanship indicators, internationalism, self-monitoring, evangelical identification, age, gender, education, income). Living standards is scaled 0–1; paternalism and agency are on 1–5 scales. HC2 robust standard errors in parentheses. Data: Prolific, August 2026.

Figure A9: Predicted probability of switching, by paternalism



Notes: Exploratory and descriptive. Predicted probability that an omnibus opponent ($n = 811$) supports at least one itemized program, across the range of the post-treatment paternalism index, from a logistic regression with the full covariate set (ideology, partisanship indicators, internationalism, self-monitoring, evangelical identification, age, gender, education, income) held at its means; HC2 standard errors. The shaded band denotes the 95% confidence interval. Data: Prolific, August 2026.

7.1 Sensitivity of the Pathway Estimates

Analysis status: Non-preregistered sensitivity analysis of the preregistered pathway models.

Table A28 reports the parametric sensitivity analysis for the living-standards pathway alongside the preregistered paternalism analysis. The living-standards estimate is overturned by relatively small residual correlations, whereas the paternalism null is more robust.

Table A28: Sensitivity of the mediation estimates to unmeasured confounding

Mediator	Outcome	ACME	95% CI	$ \rho $ at ACME = 0
Perceived living standards	Omnibus support	0.0050	[0.0008, 0.0100]	0.05
Perceived living standards	Cash-item support	0.0065	[0.0027, 0.0115]	0.10
Perceived living standards	Humanitarian support	0.0069	[0.0029, 0.0120]	0.10
Paternalism (composite)	Omnibus support	0.0049	[−0.0030, 0.0139]	0.30

Notes: Each row reports the average causal mediation effect of the Black-country treatment through the named mediator and the absolute correlation ρ between the mediator- and outcome-equation error terms at which the ACME crosses zero, from the parametric sensitivity analysis of Imai, Keele & Yamamoto (2010) as implemented in `medsens` (grid of 0.05). Outcomes are scaled 0–1; both equations include the full covariate set. The living-standards channel, the only ACME to clear conventional significance, is overturned by substantially smaller violations of sequential ignorability than the paternalism null. Main arm, $N = 1,769$. Data: Prolific, August 2026.

8 Additional Descriptive Analyses

8.1 Correlates of Conditional Support

Analysis status: Exploratory and descriptive; not preregistered. The analysis conditions on post-treatment omnibus opposition.

Table A29 regresses support for at least one program on respondent characteristics among omnibus opponents. These associations describe which observed opponents express conditional support and are not interpreted causally.

Table A29: Who switches: correlates of supporting at least one program among omnibus opponents

	(1)	(2)	(3)
Black-majority cue	-0.003 (0.029)	0.007 (0.028)	0.004 (0.028)
Ideology (conservatism)	-0.134*** (0.035)	-0.057 (0.039)	-0.060 (0.037)
Independent	-0.014 (0.053)	-0.025 (0.051)	-0.006 (0.049)
Republican	0.061 (0.073)	0.039 (0.069)	0.032 (0.066)
Internationalism	0.095*** (0.015)	0.085*** (0.015)	0.051*** (0.015)
Evangelical	0.063* (0.034)	0.043 (0.034)	0.047 (0.033)
Age	-0.003 (0.016)	-0.004 (0.015)	-0.002 (0.015)
Male	-0.015 (0.030)	0.007 (0.030)	0.014 (0.029)
Education	-0.003 (0.015)	-0.001 (0.015)	0.007 (0.015)
Income	0.025 (0.016)	0.020 (0.016)	0.019 (0.016)
Ethnocentrism		-0.095*** (0.018)	-0.088*** (0.017)
Racial resentment		-0.012 (0.025)	0.019 (0.023)
Paternalism			0.106*** (0.015)
Mean of dep. var.	0.764	0.764	0.764
Observations	811	811	811
R ²	0.099	0.141	0.190

Notes: Linear probability models of switching (funding at least one itemized program) among respondents who oppose the omnibus package ($n = 811$), with HC2 robust standard errors in parentheses. Continuous predictors (ideology, internationalism, age, education, income, ethnocentrism, racial resentment, paternalism) are standardized; party indicators use the Democratic baseline. Paternalism is the post-treatment index. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. Data: Prolific, August 2026.

8.2 Geographic References in Open-Ended Responses

Analysis status: Exploratory and descriptive; not preregistered.

Table A30: Places named in open-ended responses, by recipient-race condition

Place	White-majority (<i>n</i> = 872)	Black-majority (<i>n</i> = 897)	Total
<i>Places the respondent supplied</i>			
Africa	0	6	6
Europe (any)	3	1	4
Ukraine	3	1	4
China	0	3	3
Israel or Palestine	2	1	3
Haiti	1	1	2
Argentina	0	1	1
Iran	0	1	1
Iraq	0	1	1
Russia	1	0	1
Taiwan	0	1	1
<i>References the instrument invites</i>			
Glanton (vignette country)	82	104	186
United States	91	100	191

Notes: Exploratory; not preregistered. Entries are counts of respondents whose open-ended response names each place, by randomized condition. Places named by no respondent are omitted. The vignette country is fictional and no region is mentioned in the instrument, so every entry is a geography the respondent supplied. Counts are too small for inferential comparisons and are reported descriptively. Main arm, *N* = 1,769. Data: Prolific, August 2026.

8.3 Perceived Region and Religion of the Recipient Country

Analysis status: Exploratory and descriptive; not preregistered.

Table A31: Perceived religion and region of the recipient country, by treatment

	White-majority	Black-majority
Expected majority religion		
Christian	67.7	44.9
Muslim	18.4	42.4
Atheist	9.8	7.1
Buddhist	4.1	5.6
Expected region		
Sub-Saharan Africa	8.9	63.1
Middle East and North Africa	17.1	29.3
Latin America and Caribbean	6.6	4.2
Oceania	2.4	0.9
Asia	2.5	0.7
Central Asia	3.4	0.7
Eastern Europe	44.6	0.7
Western Europe	14.4	0.4
Observations	987	1,009

Notes: Column percentages. Respondents were asked, after the outcome and mediator batteries, what majority religion and world region they expected Glanton to have; the vignette named neither. Categories are ordered by their share in the Black-majority condition. Differences across conditions are large for both inferences (religion: $\chi^2 = 146.6$, $df = 3$, $p < .001$; region: $\chi^2 = 1039.8$, $df = 7$, $p < .001$). Includes all treated respondents surviving the registered exclusions. Data: Prolific, August 2026.

9 Survey Instrument

Question wording below is transcribed from the fielded instrument. Items appear in the order respondents encountered them. Response options are given in brackets after each item. Program order was randomized at the respondent level and carried through every subsequent battery.

9.1 Pre-Treatment Demographics

1. In what year were you born?
2. What is your sex?
3. What is the highest level of education you have completed?
4. Income
5. Would you describe yourself as a “born-again” or evangelical Christian?
6. Which of the following best describes your race or ethnicity?
7. Are you of Spanish, Latino, or Hispanic origin or descent?

9.2 Political Predispositions

1. Generally speaking, do you usually think of yourself as a Democrat, a Republican, an Independent, or what? [Followed by the standard strength and leaner probes.]
2. In general, do you think of yourself as: [Extremely liberal; Liberal; Slightly liberal; Moderate, middle of the road; Slightly conservative; Conservative; Extremely conservative]
3. The U.S. needs to play an active role in world affairs. (*Internationalism*) [Agree strongly–Disagree strongly]

Ethnocentrism

Six items, five-point agree-disagree. Items 4-6 are reverse coded. $\alpha = .77$.

1. No matter what happens, I will ALWAYS support my ethnic group and never let it down.
2. It is absolutely vital that all members of our ethnic group think and behave as one.
3. In general, I prefer doing things with people from my own ethnic group rather than with people from other ethnic groups.
4. I don't believe that my ethnic group is any better than any other. (*reversed*)
5. I like the idea of a society in which people from completely different ethnic groups mix together freely. (*reversed*)
6. We should always be considerate for the welfare of people from other ethnic groups even if, by doing this, we may lose some advantage over them. (*reversed*)

Racial Resentment

Four ANES items, five-point agree-disagree. Items 1 and 3 are reverse coded. $\alpha = .94$.

1. Generations of slavery and discrimination have created conditions that make it difficult for Black Americans to work their way out of the lower class. (*reversed*)
2. It's really a matter of some people not trying hard enough; if Black Americans would only try harder they could be just as well off as White Americans.
3. Over the past few years, Black Americans have gotten less than they deserve. (*reversed*)
4. Irish, Italians, Jewish and many other minorities overcame prejudice and worked their way up. Black Americans should do the same without any special favors.

Self-Monitoring


1. When you are with other people, how often do you put on a show to impress or entertain them? [Never; Once in a while; Some of the time; Most of the time; Always]
2. When you are in a group of people, how often are you the center of attention? [same scale]
3. How good or poor of an actor would you be? [Poor; Fair; Good; Excellent]

Attention Checks

Respondents failing both were screened out before treatment.

1. We would like to get a sense of your consumption of political news. To demonstrate that you've read this much, select rarely from the options below, no matter how often you watch political news. Based on the text you read above, how often do you watch political news on TV?
2. Please click the "neither agree nor disagree" response.

9.3 Treatment



The Leader of Glantonía

The United States is considering providing \$50 million in foreign aid to Glantonía.

The people in this country earn around \$5,000 a year. The largest racial group is black [white].

9.4 Manipulation Check

1. What is the largest racial group in Glantonía? [Black; White; Asian; Pacific Islander; Arab]
2. What was the name of the hypothetical country mentioned in the prompt? [Glantonía, Orevia, Tavola, Yorkshire]

9.5 Primary Outcomes

Main-arm respondents received the omnibus item followed by the itemized checklist; the itemized-only arm received the checklist alone.

1. **Omnibus.** The United States is considering providing \$50 million in foreign aid to Glantonía. Foreign aid can be used for various programs in Glantonía, such as medicine, food, military equipment, democratic elections, cash to the government or citizens, and

transportation infrastructure. Considering this would result in a modest increase in taxes, would you support or oppose foreign aid for Glantonía? [Support; Support somewhat; Neither support nor oppose; Oppose somewhat; Oppose]

2. **Itemized checklist.** [Same preamble.] Considering this would result in a modest increase in taxes, please select each program that you support funding for Glantonía. [Fund; Do Not Fund — asked separately for: Medicine; Food; Military equipment; Democratic elections; Transportation infrastructure; Cash to the government; Cash to the citizens]
3. **Allocation.** Imagine that you are in charge of allocating the \$50 million for Glantonía's programs. Please indicate the percentage of the \$50 million you would allocate to each program. The percentages should total 100%. [Constrained-sum sliders for the same seven programs, with a pie chart displaying the running allocation]

9.6 Delivery Mode

There are many ways to assist. Each type of aid can be provided as training, cash, or supplies. Please rank your preference for how to provide each assistance program. 1 = the best, 2 = middle, and 3 = the worst.

Asked for the five non-cash programs (medicine, food, military equipment, transportation infrastructure, democratic elections), each ranked across training, cash, and supplies.

9.7 Open-Ended Item

Please write a few sentences about what factors you were considering when assessing U.S. assistance to the hypothetical country of Glantonía. Please refrain from including personal information.

9.8 Post-Treatment Mediators

Presented in randomized order after a reminder of the treatment condition. All are five-point bipolar items anchored by the statements shown, except where noted.

Paternalism: Benevolence ($\alpha = .70$)

1. Please indicate which of the two statements is closer to your views: [1 = Poor countries' citizens can determine what's best for their own well-being without American guidance; 5 = Poor countries' citizens need American guidance to improve their well-being]
2. Please indicate which of the two statements is closer to your views: [1 = Poor foreigners do not need the U.S. government's protection; 5 = The U.S. government should protect poor foreigners]

Paternalism: Control ($\alpha = .69$)

1. Please indicate which of the two statements is closer to your views: [1 = The U.S. government should respect poor governments' decisions, even ones it disagrees with; 5 = The U.S. government should oppose poor governments' decisions if it believes doing so would help the citizens]
2. Please indicate which of the two statements is closer to your views: [1 = The U.S. government should respect poor governments' policies as reflecting local values, even ones it views as harmful; 5 = The U.S. government should intervene when poor governments create policies the U.S. views as harmful to their citizens]

Perceived Agency (Baker 2015; $\alpha = .74$)

1. Please indicate which of the two statements is closer to your views: [1 = Because of difficult economic circumstances, people in poor countries are unable to help themselves get richer; 5 = People in poor countries are able to do things that can help themselves

get richer]

2. Please indicate which of the two statements is closer to your views: [1 = There is a lot that poor countries and their citizens can do by themselves to improve their own livelihoods; 5 = There is little that poor countries and their citizens can do by themselves to improve their own livelihoods] (*reversed*)
3. Please indicate the extent to which you agree or disagree with the following statement: When it comes to improving their economic standard of living, people in poor countries are like extremely sick or paralyzed patients; they are completely unable to help themselves. (*reversed; five-point agree-disagree*)

Perceived Living Standards (Baker 2015; $\alpha = .78$)

1. Giving your best guess, what percentage of Glantonía's households have indoor plumbing? [0–100]
2. Giving your best guess, what percentage of Glantonía's households have at least two meals per day? [0–100]

9.9 Leader and Government Perceptions

1. How qualified do you think Glantonía's leader is compared to other world leaders? [Extremely unqualified; Somewhat unqualified; Neither qualified nor unqualified; Somewhat qualified; Extremely qualified]
2. How much do you trust the government of Glantonía to use U.S. aid effectively? [Not at all; A little; Somewhat; A lot; A great deal]
3. How much do you trust the government of Glantonía to use U.S. aid honestly? [same scale]
4. How widespread do you think corruption is among government officials in Glantonía?

[Not at all widespread; A little widespread; Somewhat widespread; Very widespread; Extremely widespread]

9.10 Mechanism Batteries

Each respondent rated all seven programs on each dimension (Yes/No).

1. Please indicate if you believe these programs reduce theft by politicians in Glantonía.
2. Please indicate if you believe these programs make it easy to verify aid was used as intended by Glantonía.
3. Please indicate if you believe that these programs benefit the American economy.
4. Please indicate if you believe that these programs increase American global power.
5. Please indicate if you believe that these programs save lives.
6. Please indicate if you believe that these programs are the moral thing to do.

9.11 Inference Checks

1. What would you expect Glantonía's majority religion is? [Christian; Muslim; Buddhist; Atheist]
2. What region would you expect Glantonía to be in? [Sub-Saharan Africa; Middle East and North Africa; Western Europe; Eastern Europe; Latin America and Caribbean; Asia; Central Asia; Oceania]